

# backend

cohort #0 by open camp

# #5 's agenda

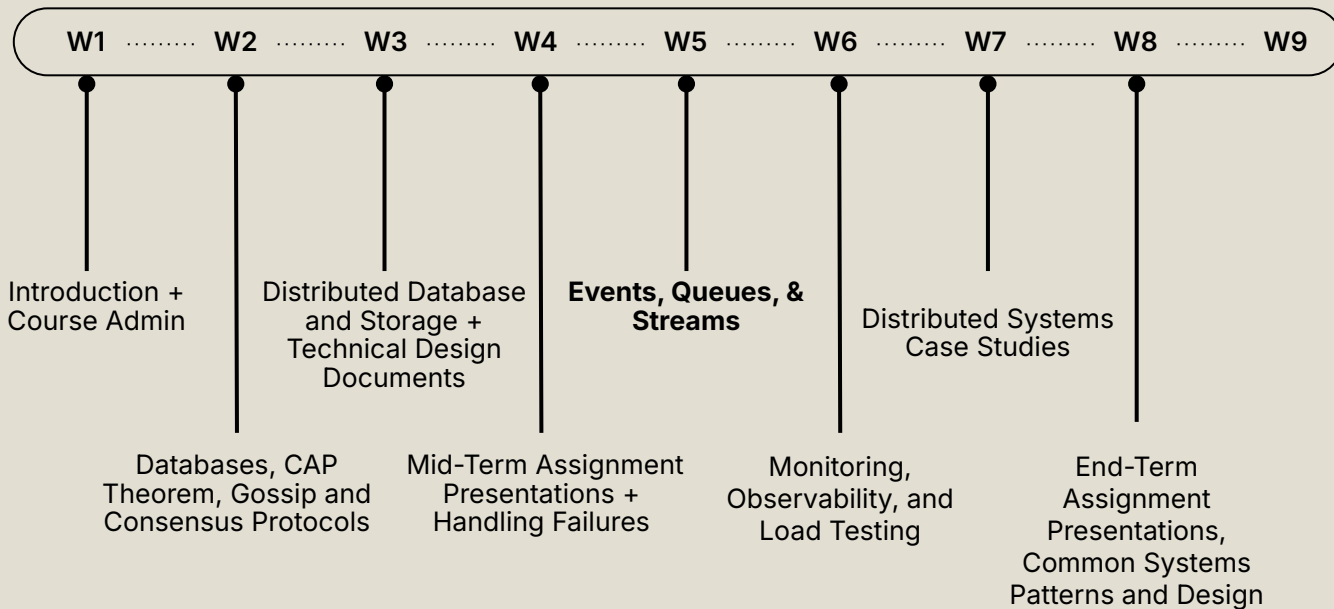
- 1 ..... admin matters (if any)
- 2 ..... all about events
- 3 ..... rabbitmq and kafka
- 4 ..... w5 assignment

**admin matters**

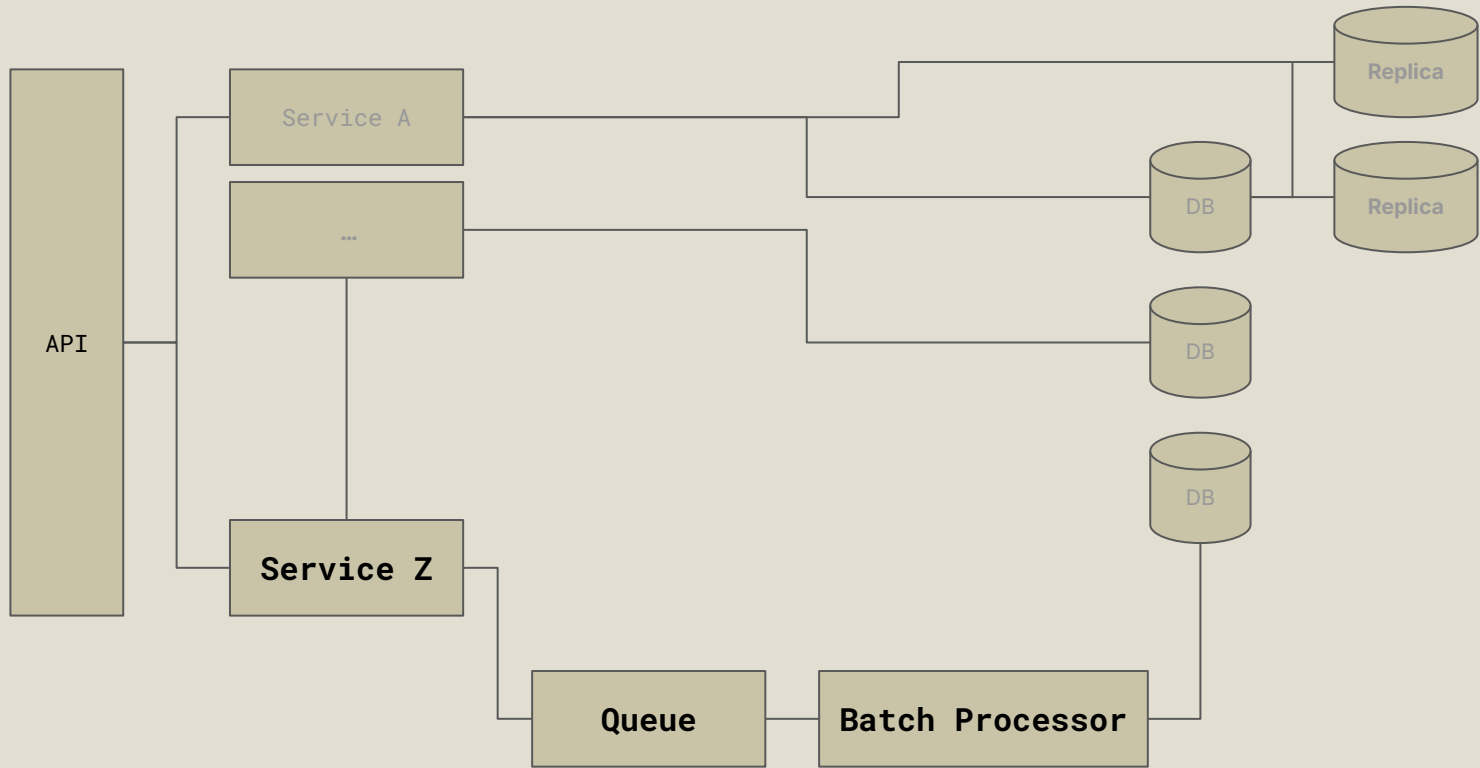
Curriculum: <https://opencamp-cc.github.io/backend-curriculum/>

Start

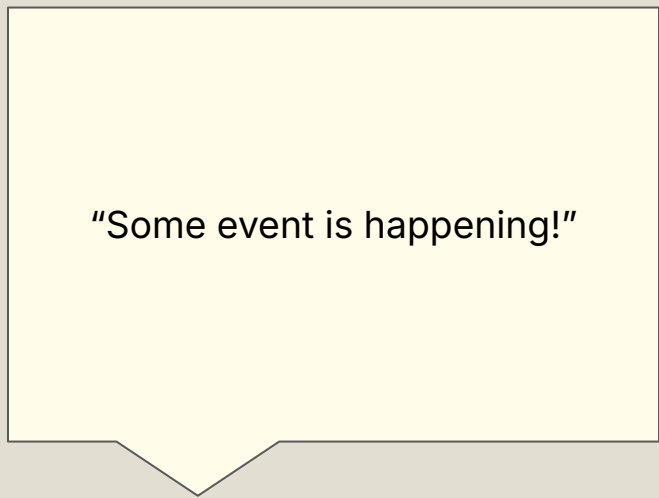
End



# 5.1 all about events



# Queues: Async Processing and Retrieval



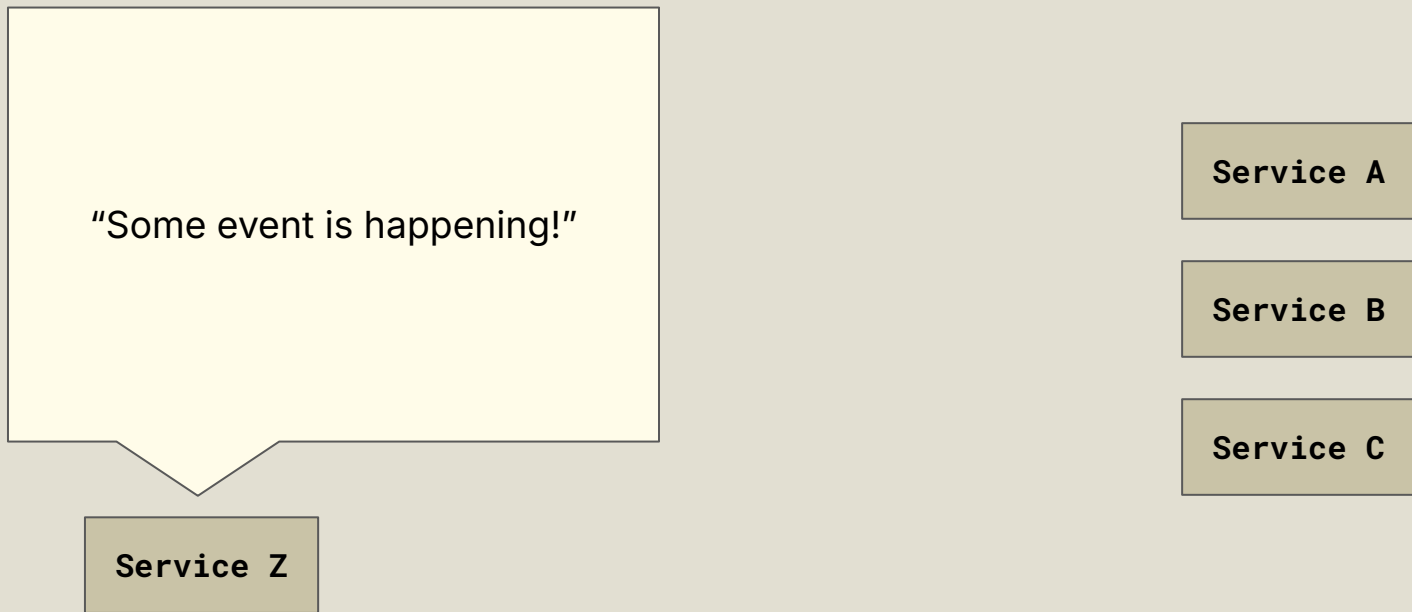
**Service Z**

**Service A**

**Service B**

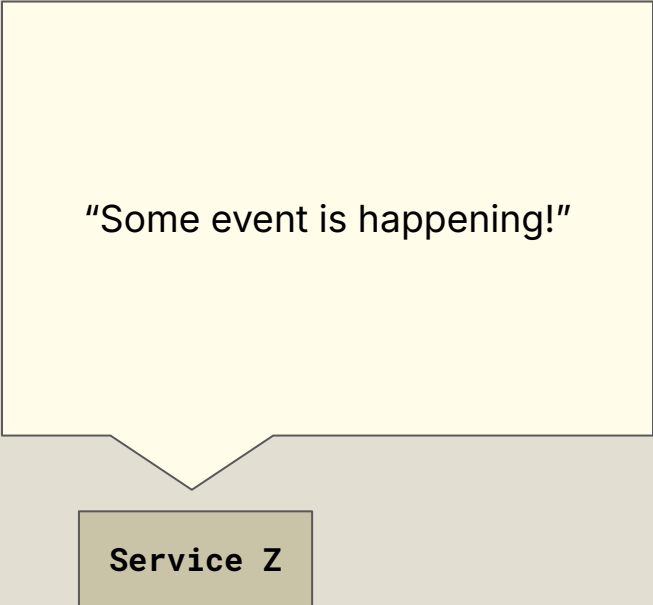
**Service C**

**Fires an event...**



**... but doesn't care if others receive it**

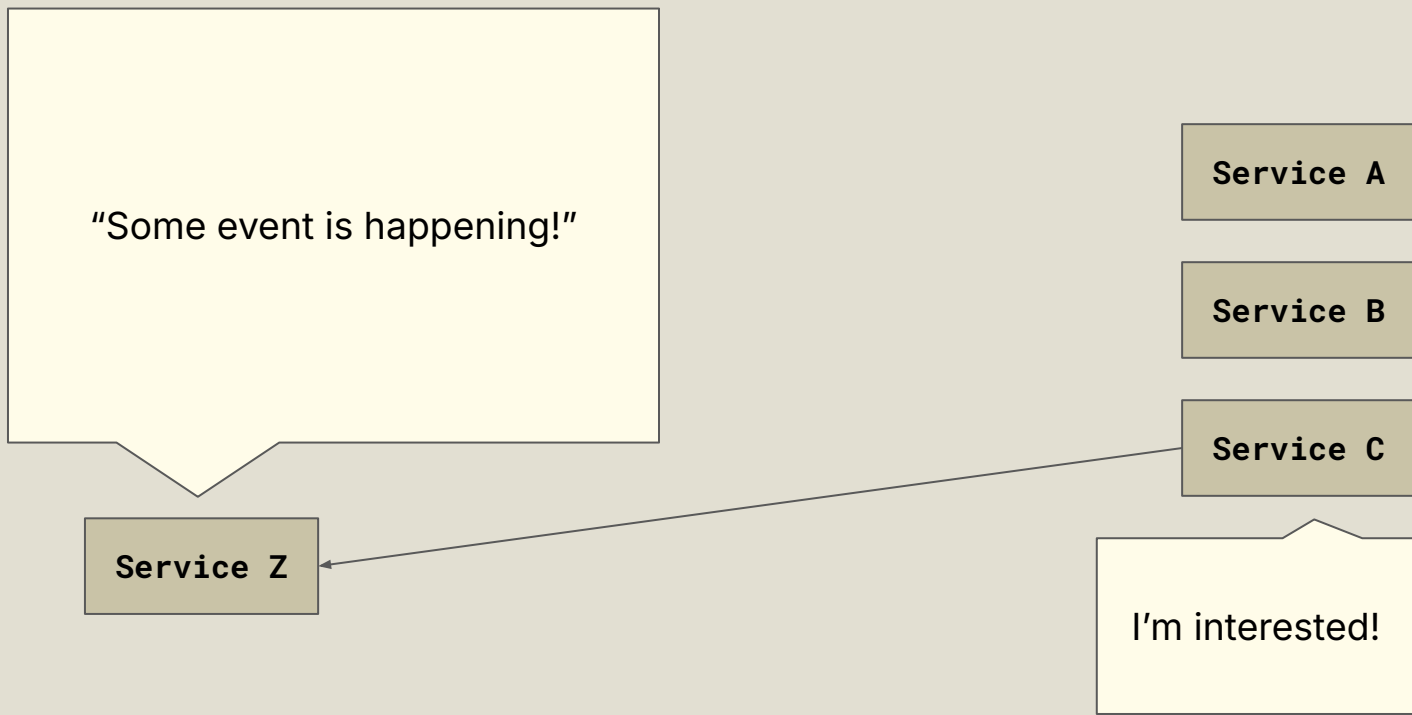




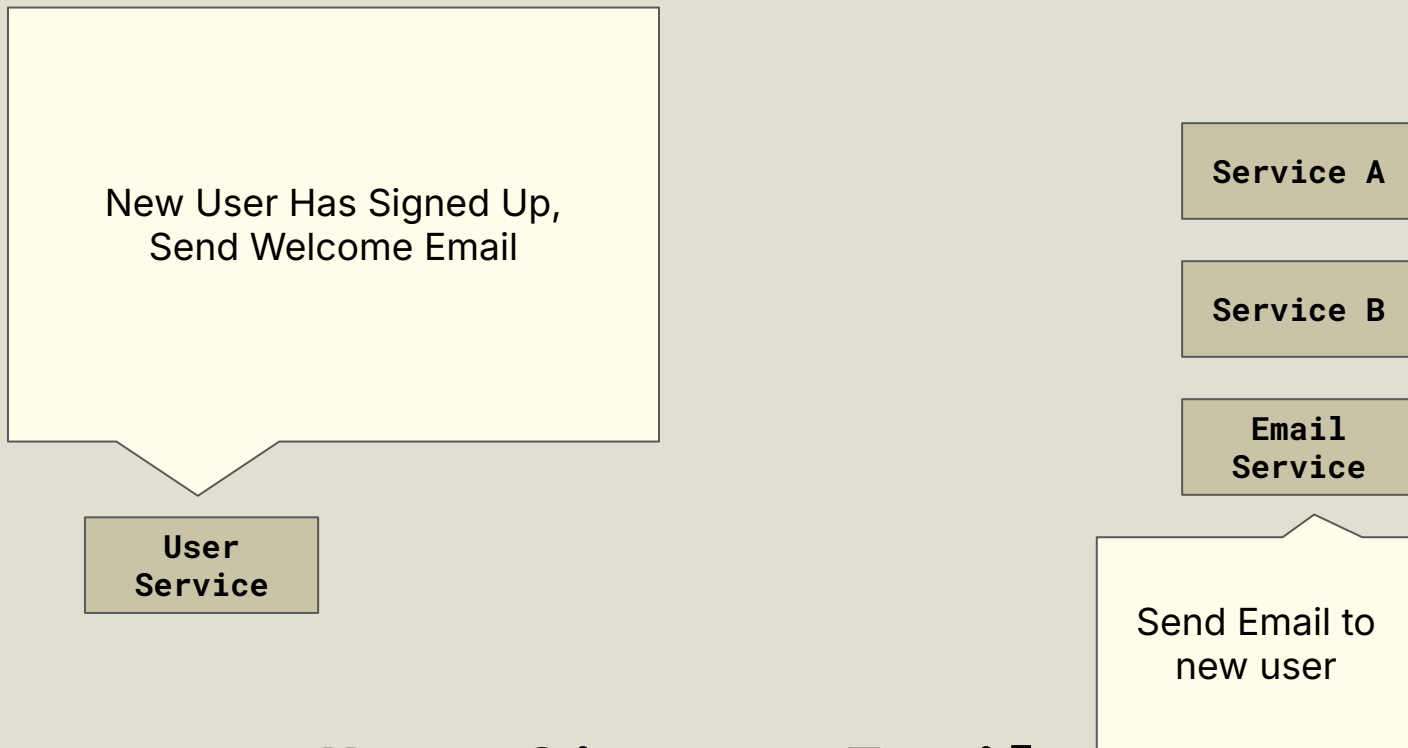
"Some event is happening!"

**Service Z**

**Fires an event...**




**... and others follow up**



# User Signup: Email

1. Pushes event

2. Receives event



**mailchimp**

# User Signup: Email

## Related to items you've viewed [See more](#)



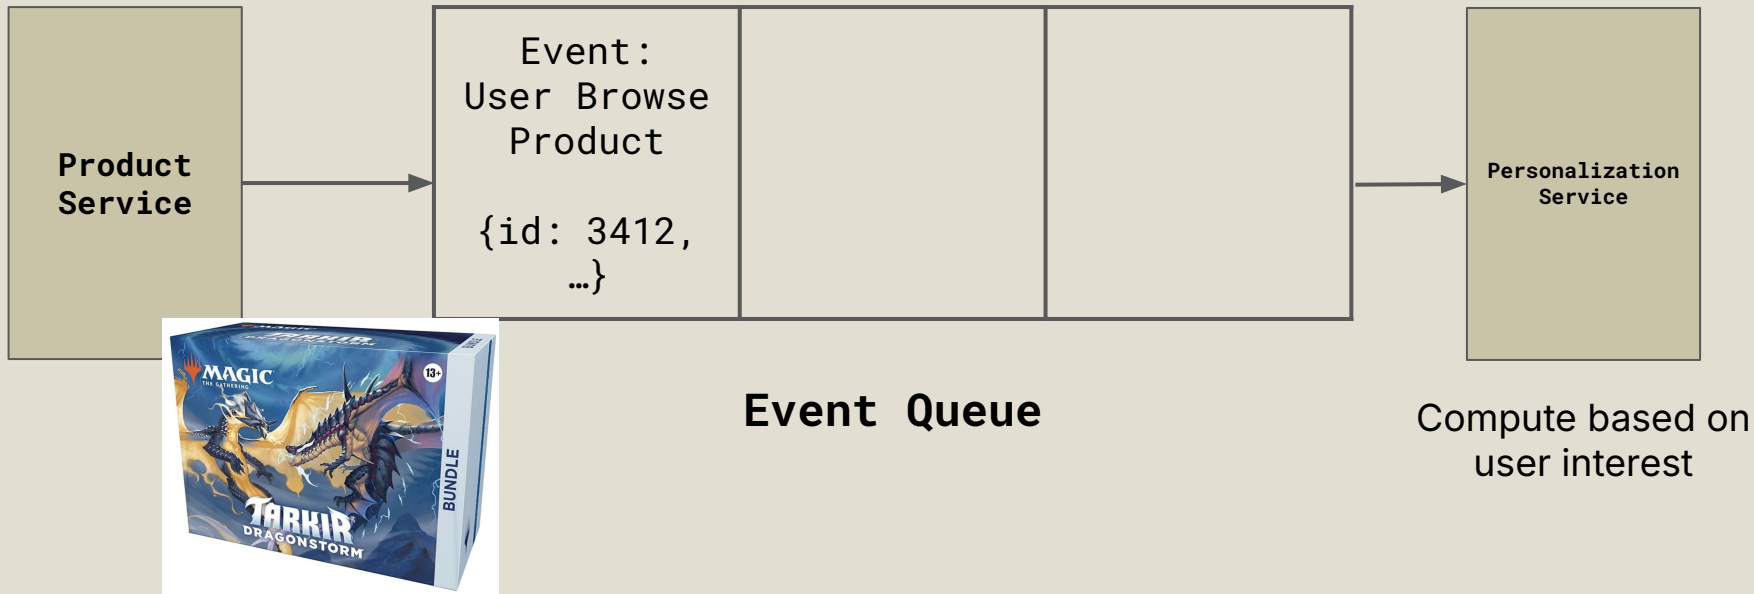
## Related to items you've viewed [See more](#)



# Personalization?

1. Pushes event

2. Receives event



# Personalization: Background Processing

Events can be:

- **Lossy**: it is possible that they get "lost" or "ignored"
- **One-way**: you are not guaranteed to get a response
- **One-to-Many**: an event can be received by one or multiple listeners (subscribers)
- **Asynchronous**: events may not be processed immediately



**Distributed Message Broker**



**Event Streaming Platform**

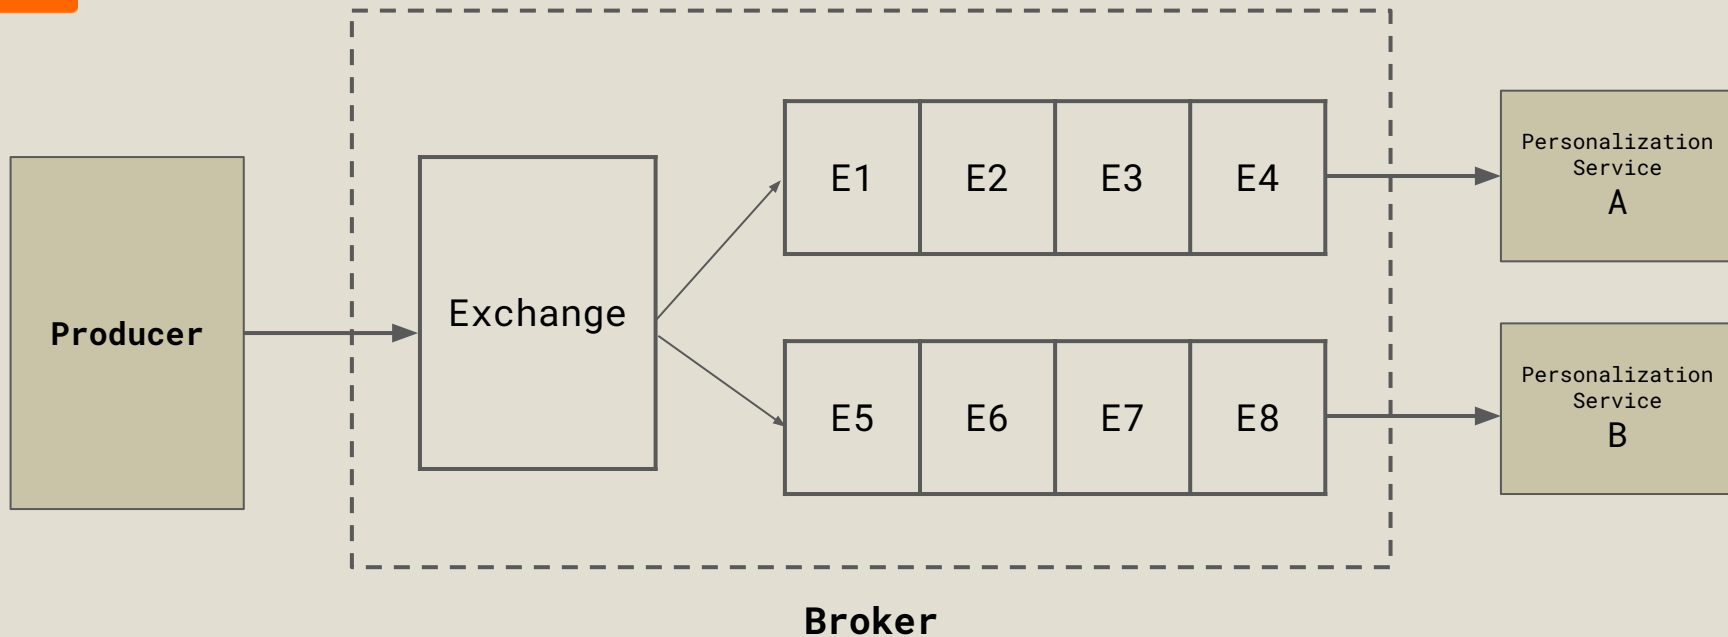
**Two Message Queue Systems**



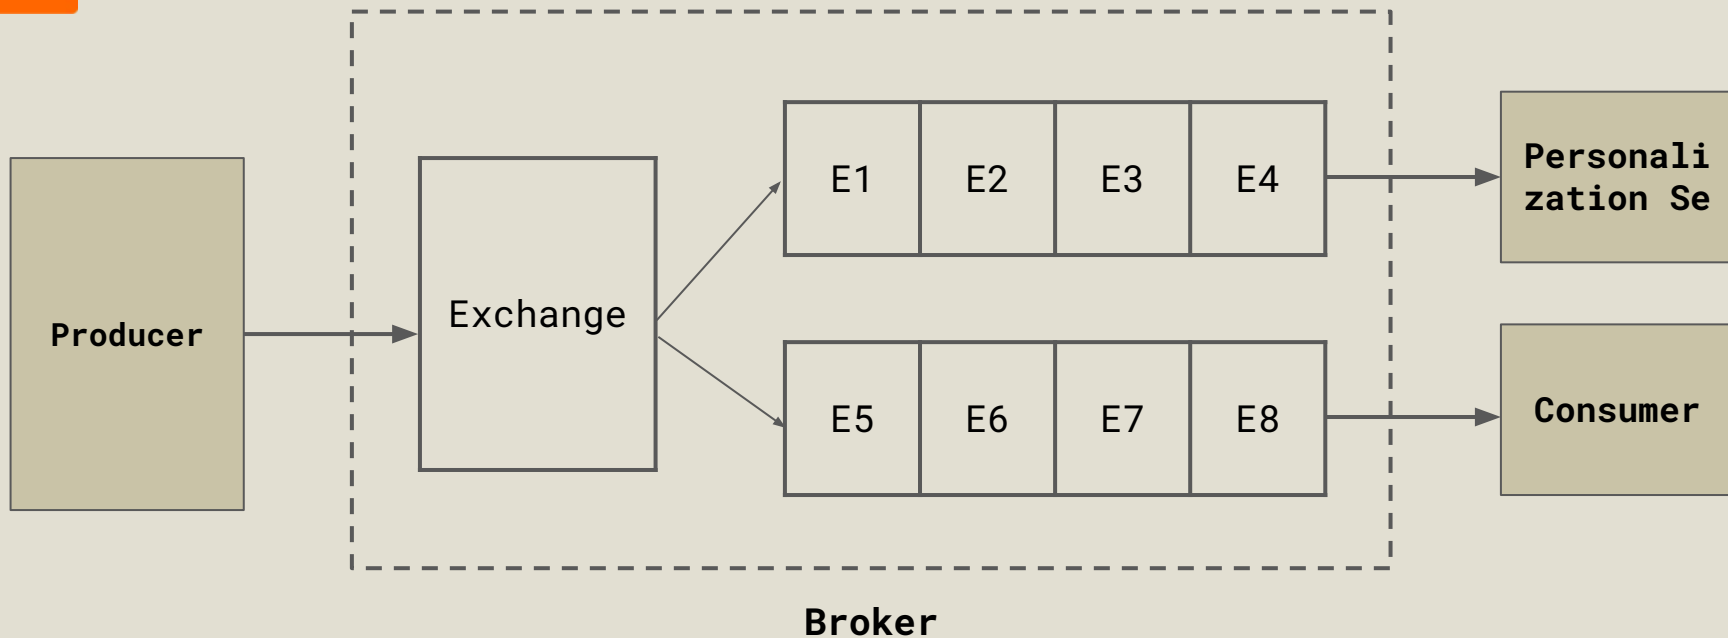
**5.1.1 rabbitmq**



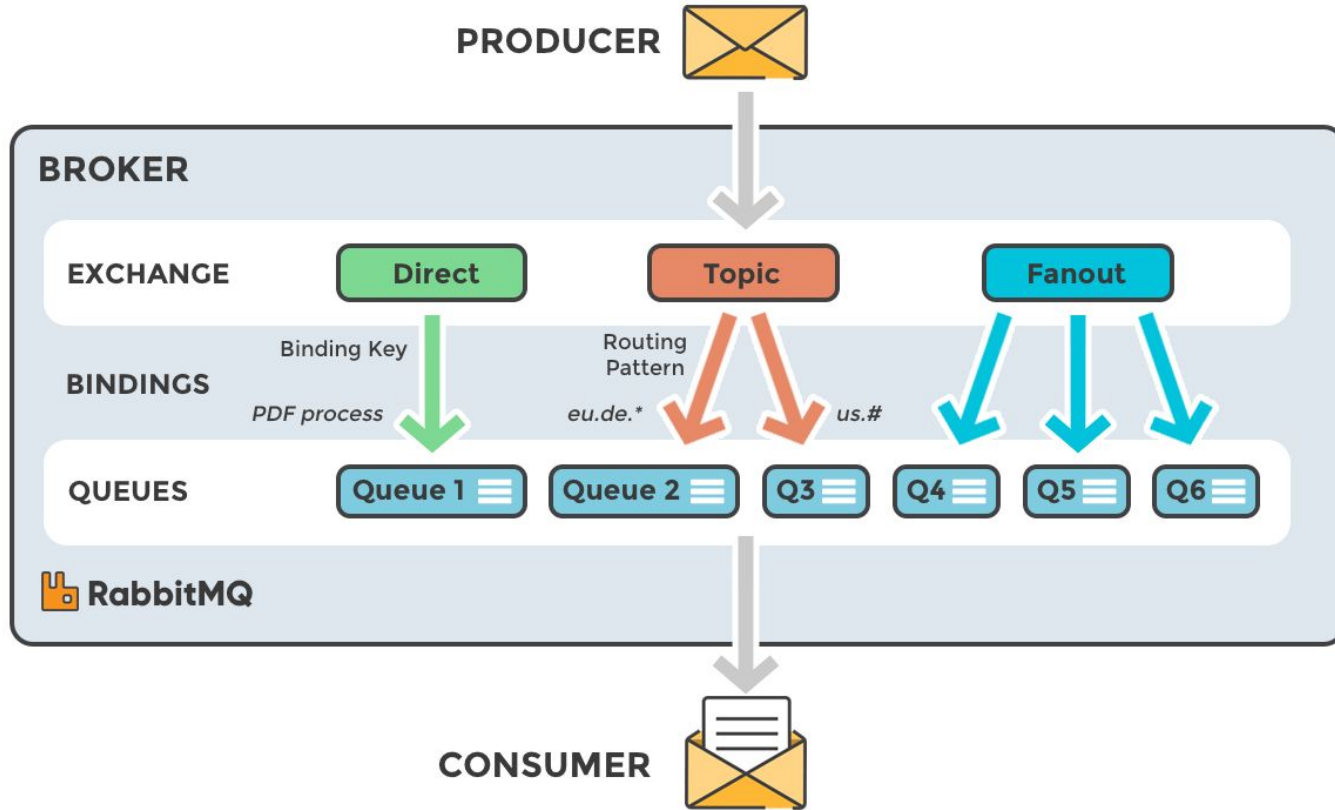
# RabbitMQ: Exchange Broker



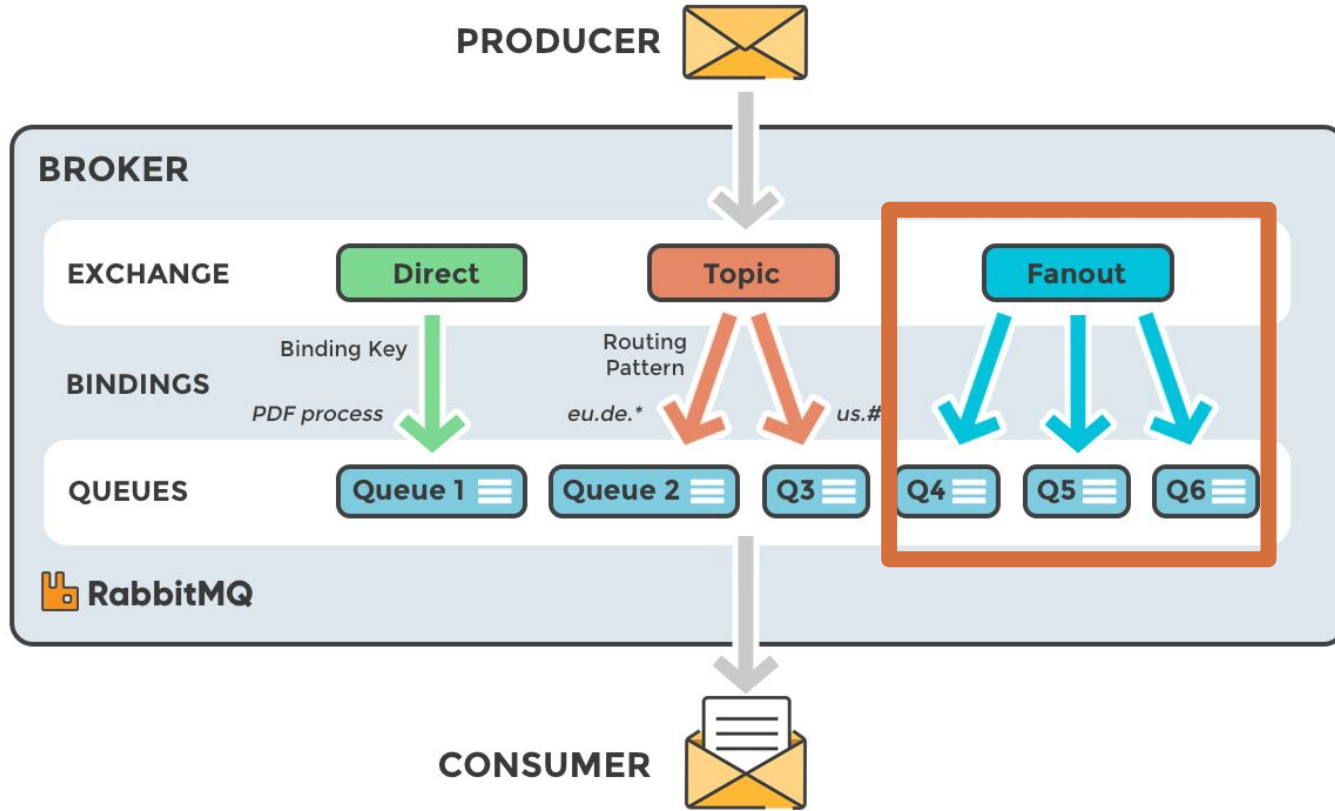
Events are “sorted” into queues at the Exchange




Events are “sorted” into queues at the Exchange



Direct, Topic, and Fanout Exchanges



**Direct, Topic, and Fanout Exchanges**



**Fanout: Event is sent to all queues**

**Related to items you've viewed** [See more](#)



**Pick up where you left off**



Magic: The Gathering... Magic: The Gathering...

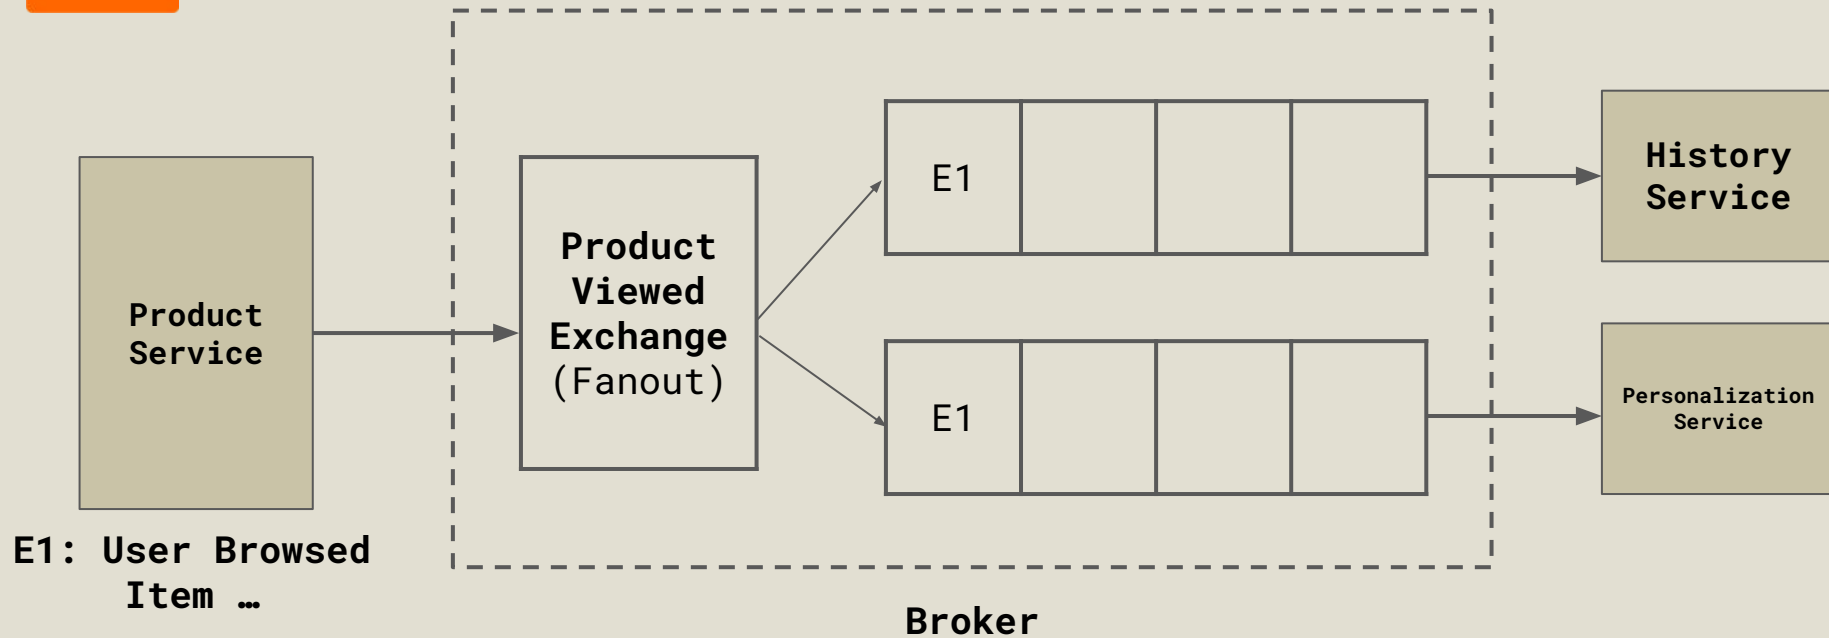


Magic: The Gathering... Magic The Gathering...

[See more](#)

**Personalization + History**





**Using Fanout for all interested services**



Product  
Service

```
channel.exchange_declare(exchange='productviews',  
                          exchange_type='fanout')
```

**Declaring (creating) an Exchange on Producer**



**productviews Exchange Created**




Personalization  
Service  
A

```
result = channel.queue_declare(queue='a_queue',  
                               exclusive=True)
```

**Declaring (creating) an Queue on Consumer**



## Bind the queues to the Exchange via Routing Key



**Queue is declared (created)**




Personalization  
Service  
**A**

```
channel.queue_bind(exchange='productviews',  
                   queue='a_queue')
```

**Binding a Queue to an Exchange (Service A)**



## Bind the queues to the Exchange via Routing Key



**Binding 'a\_queue' to 'productviews'**



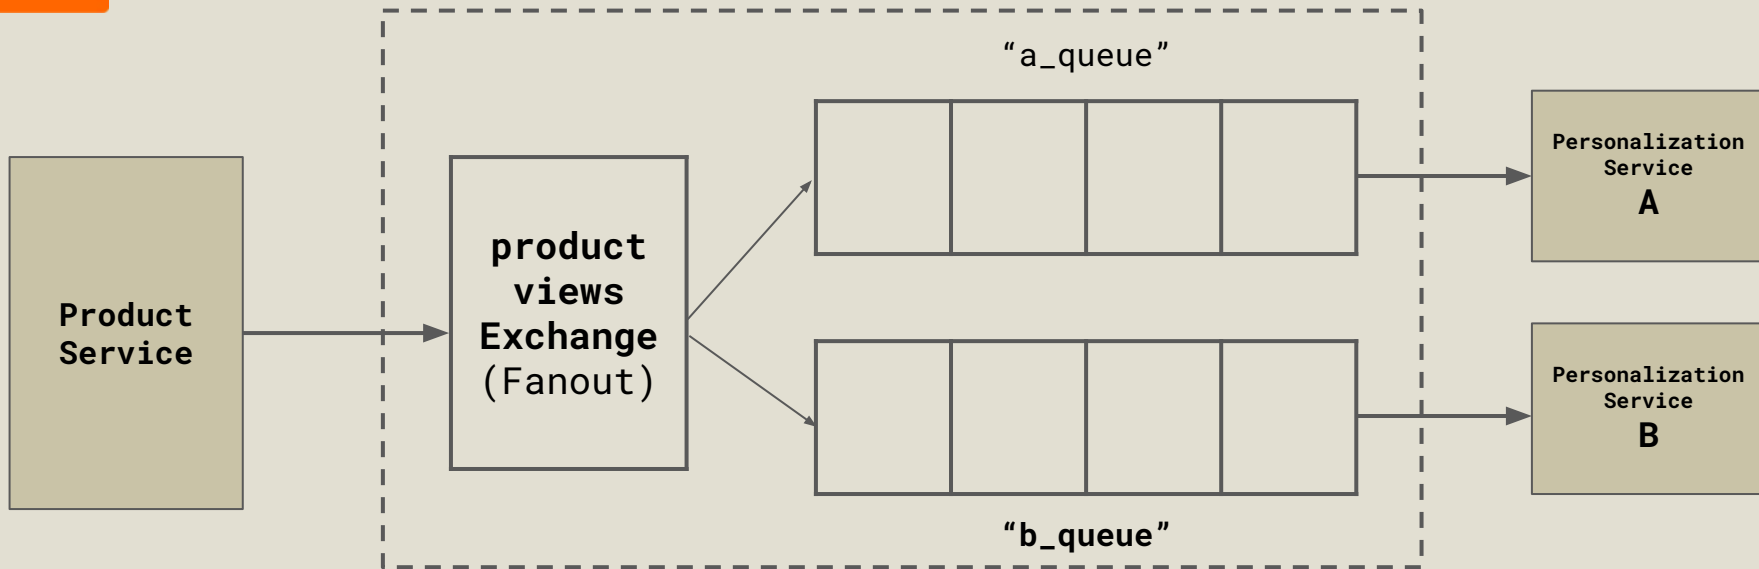
Personalization  
Service  
**B**

```
result = channel.queue_declare(queue='b_queue',  
                                exclusive=True)
```

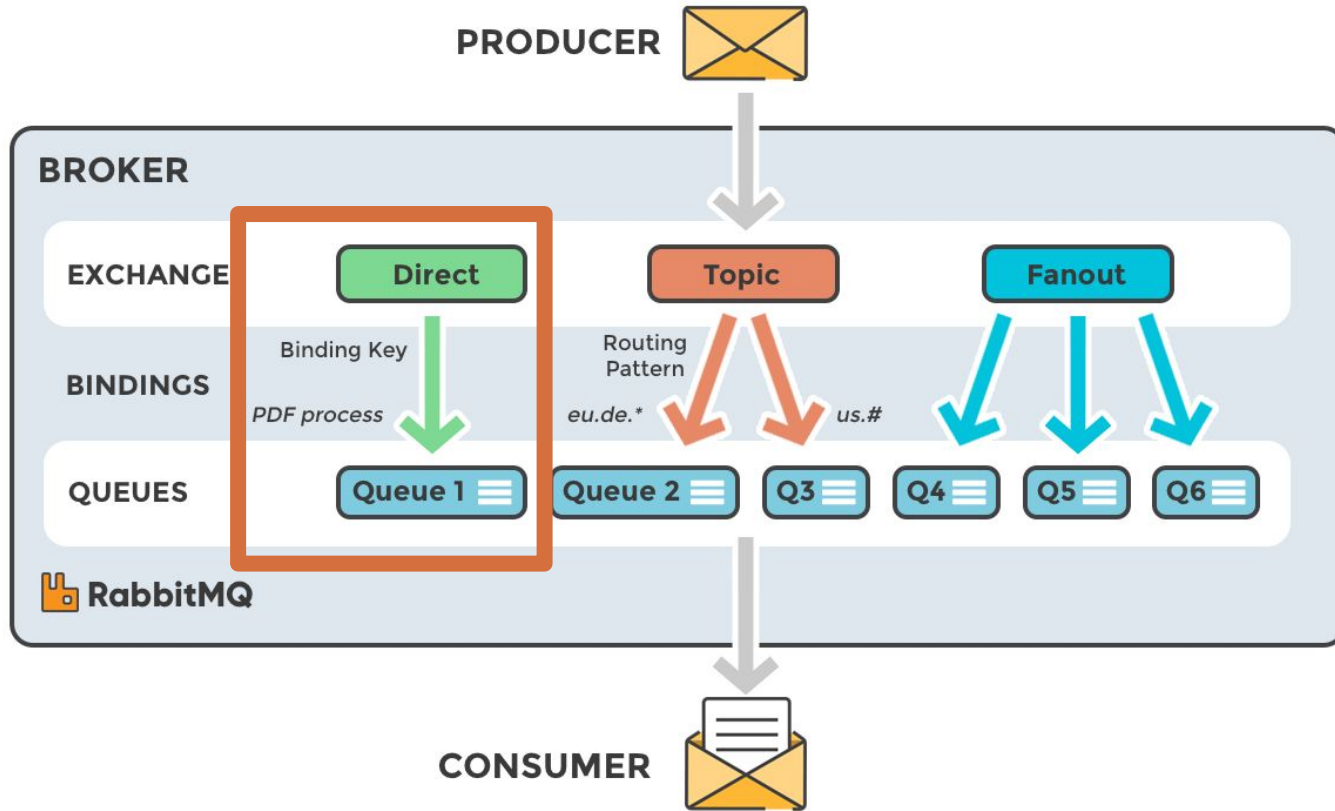
```
channel.queue_bind(exchange='productviews',  
                   queue='b_queue')
```

**Binding a Queue to an Exchange (Service B)**

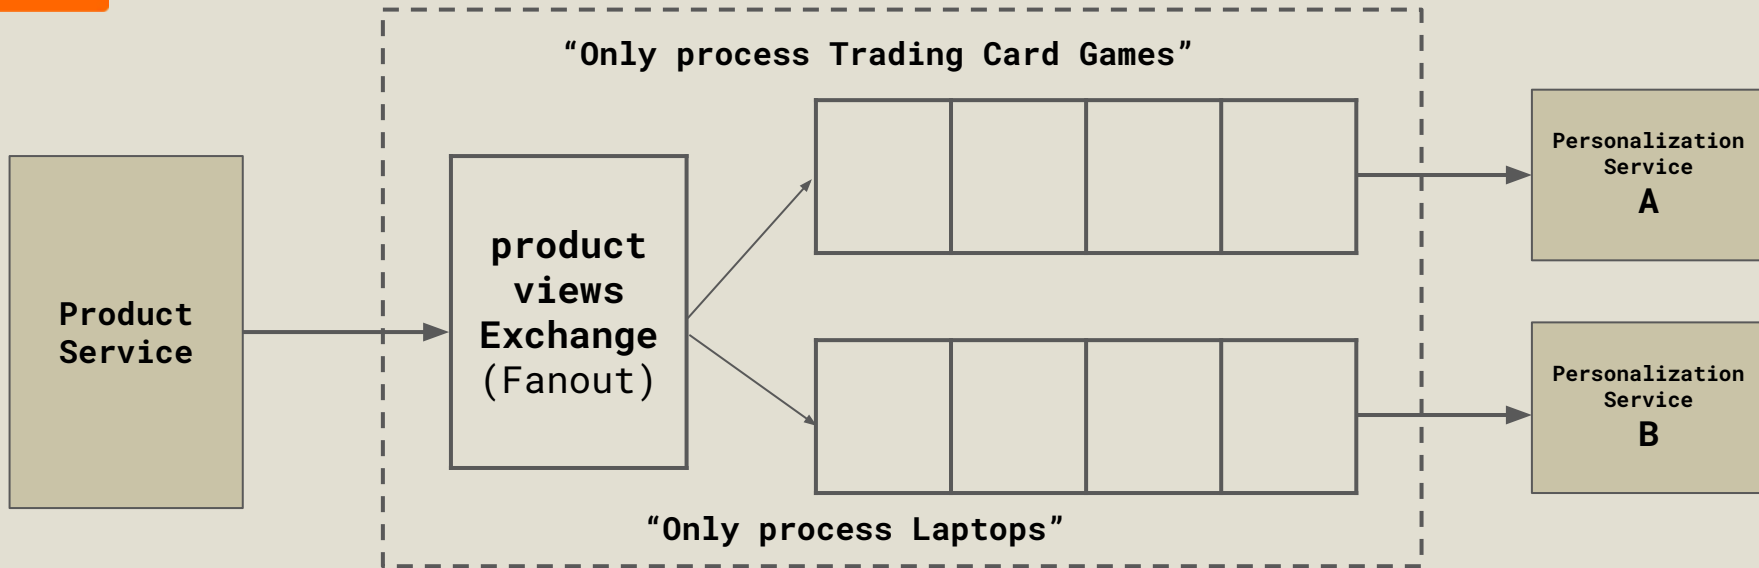




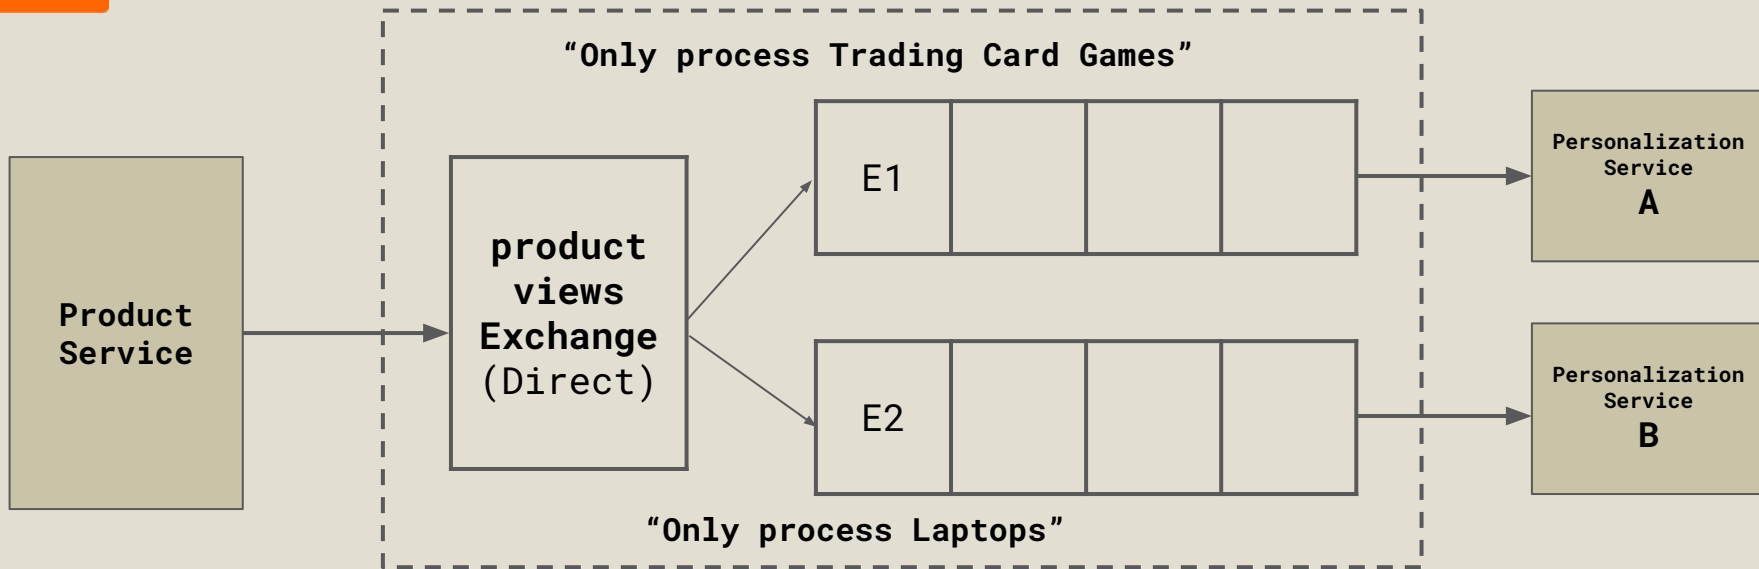
**Fanout Exchange: Bound to Two Queues**



**Direct, Topic, and Fanout Exchanges**



**Can we filter for only certain events?**



E1: Browsed Magic Cards

E2: Browsed Macbook

**Use Direct Exchange to Filter by Routing Key**



Product  
Service

```
channel.exchange_declare(exchange='productviews',  
                          exchange_type='direct')
```

**Declaring (creating) an Exchange on Producer**



**productviews Exchange Created**




Personalization  
Service  
A

```
result = channel.queue_declare(queue='a_queue',  
                               exclusive=True)
```

**Declaring (creating) an Queue on Consumer**



## Bind the queues to the Exchange via Routing Key



**Queue is declared (created)**





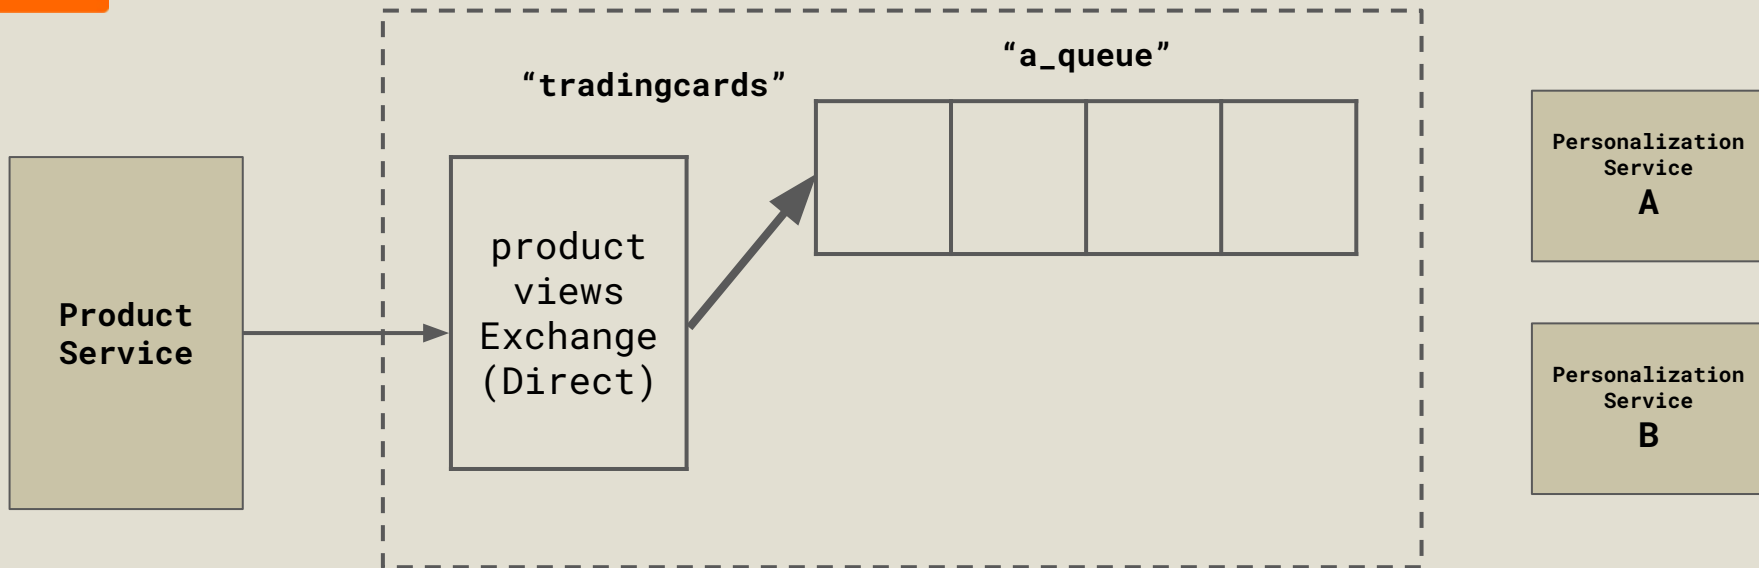
Personalization  
Service  
**A**

```
channel.queue_bind(exchange='productviews',  
                    queue='a_queue',  
                    routing_key='tradingcards')
```

**Binding a Queue to an Exchange (Service A)**



## Bind the queues to the Exchange via Routing Key



**Binding 'a\_queue' to 'productviews'**

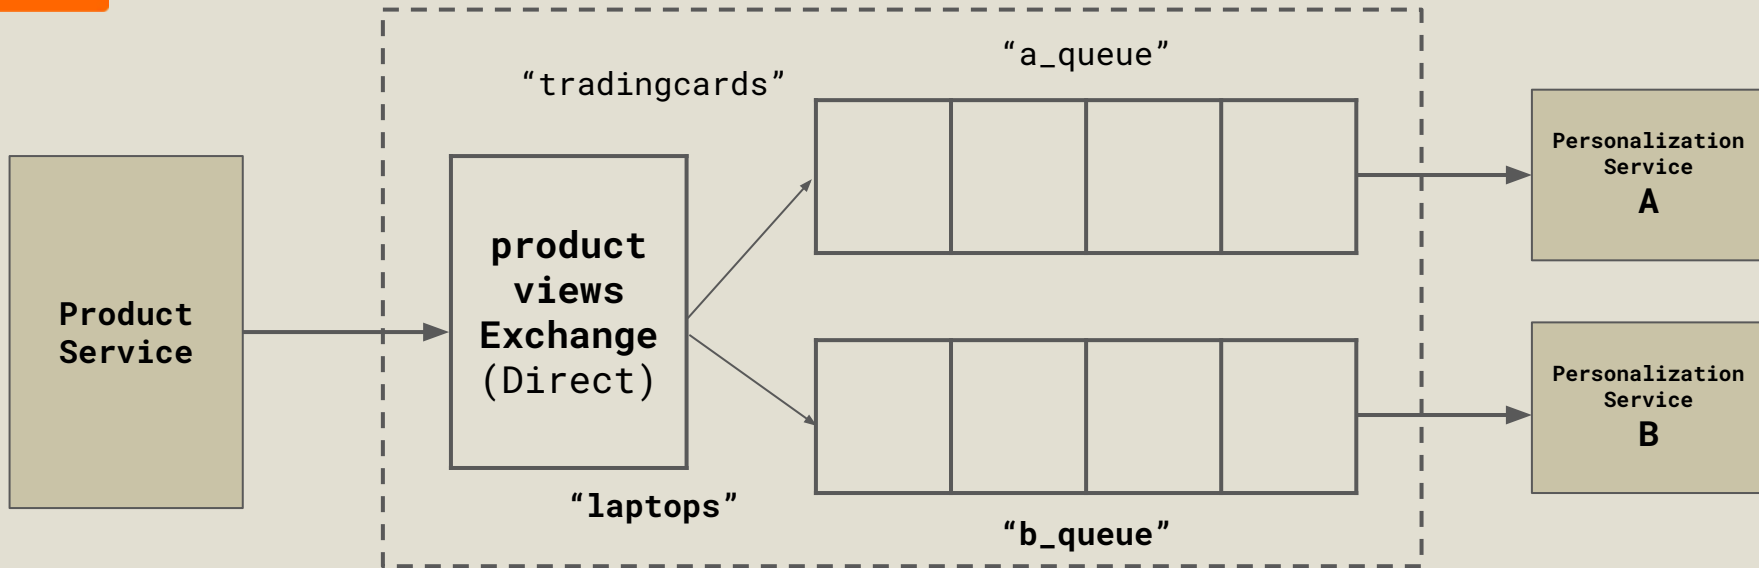


Personalization  
Service  
**B**


```
result = channel.queue_declare(queue='b_queue',  
                                exclusive=True)
```

```
channel.queue_bind(exchange='productviews',  
                   queue='b_queue',  
                   routing_key='laptops')
```

**Binding a Queue to an Exchange (Service B)**



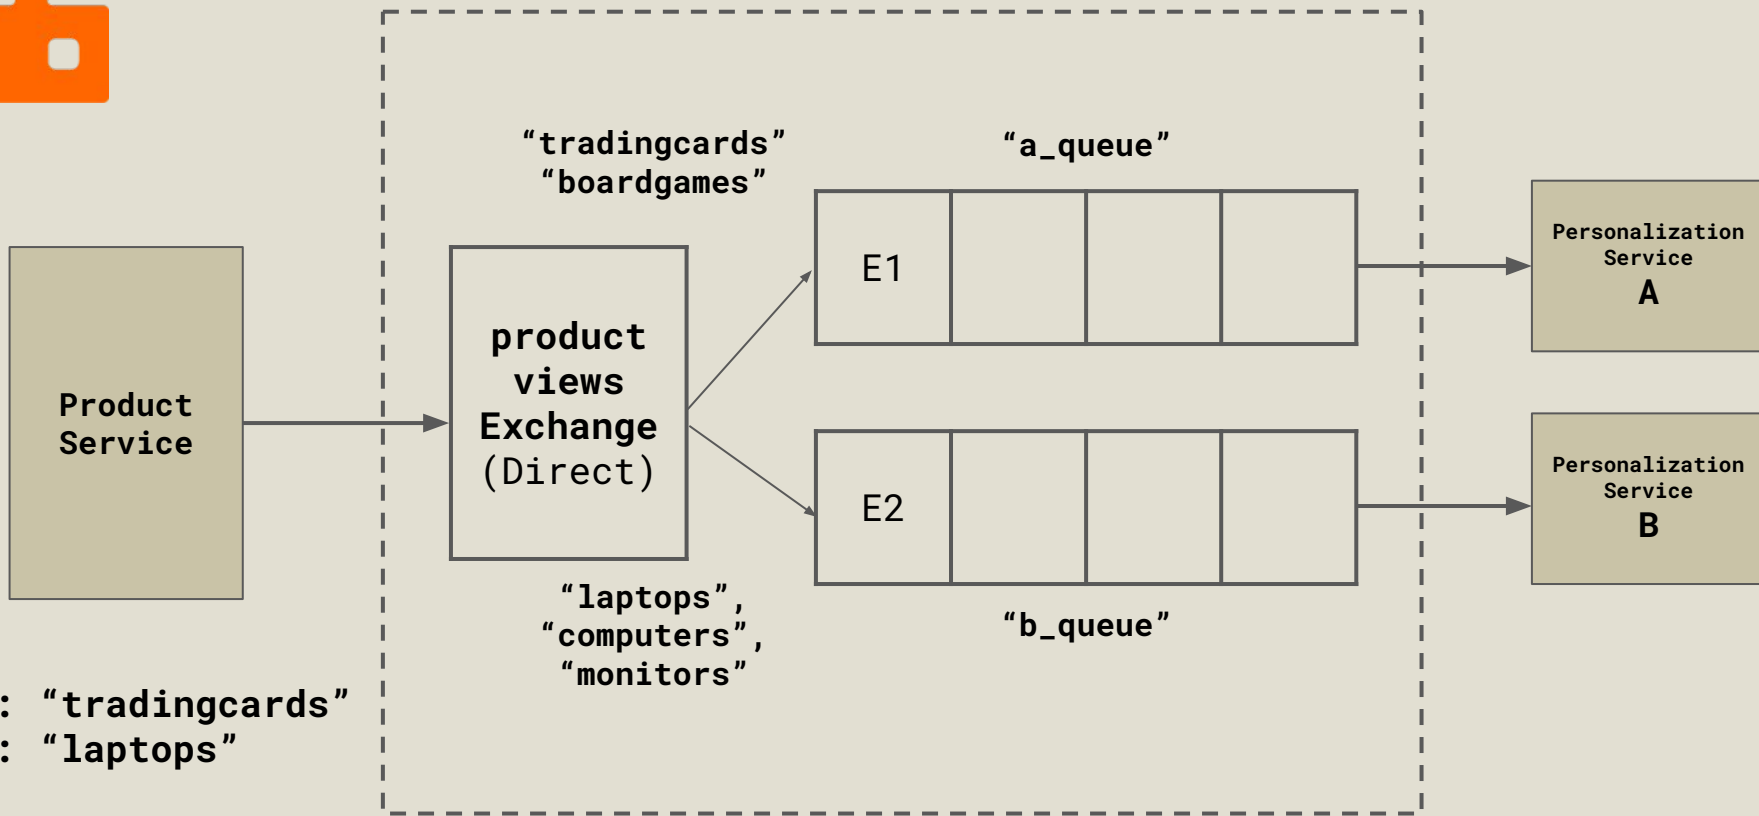
**Direct Exchange: Bind Each Category to a Queue**



E1: "tradingcards"

E2: "laptops"

**Direct Exchange: Bind Each Category to a Queue**



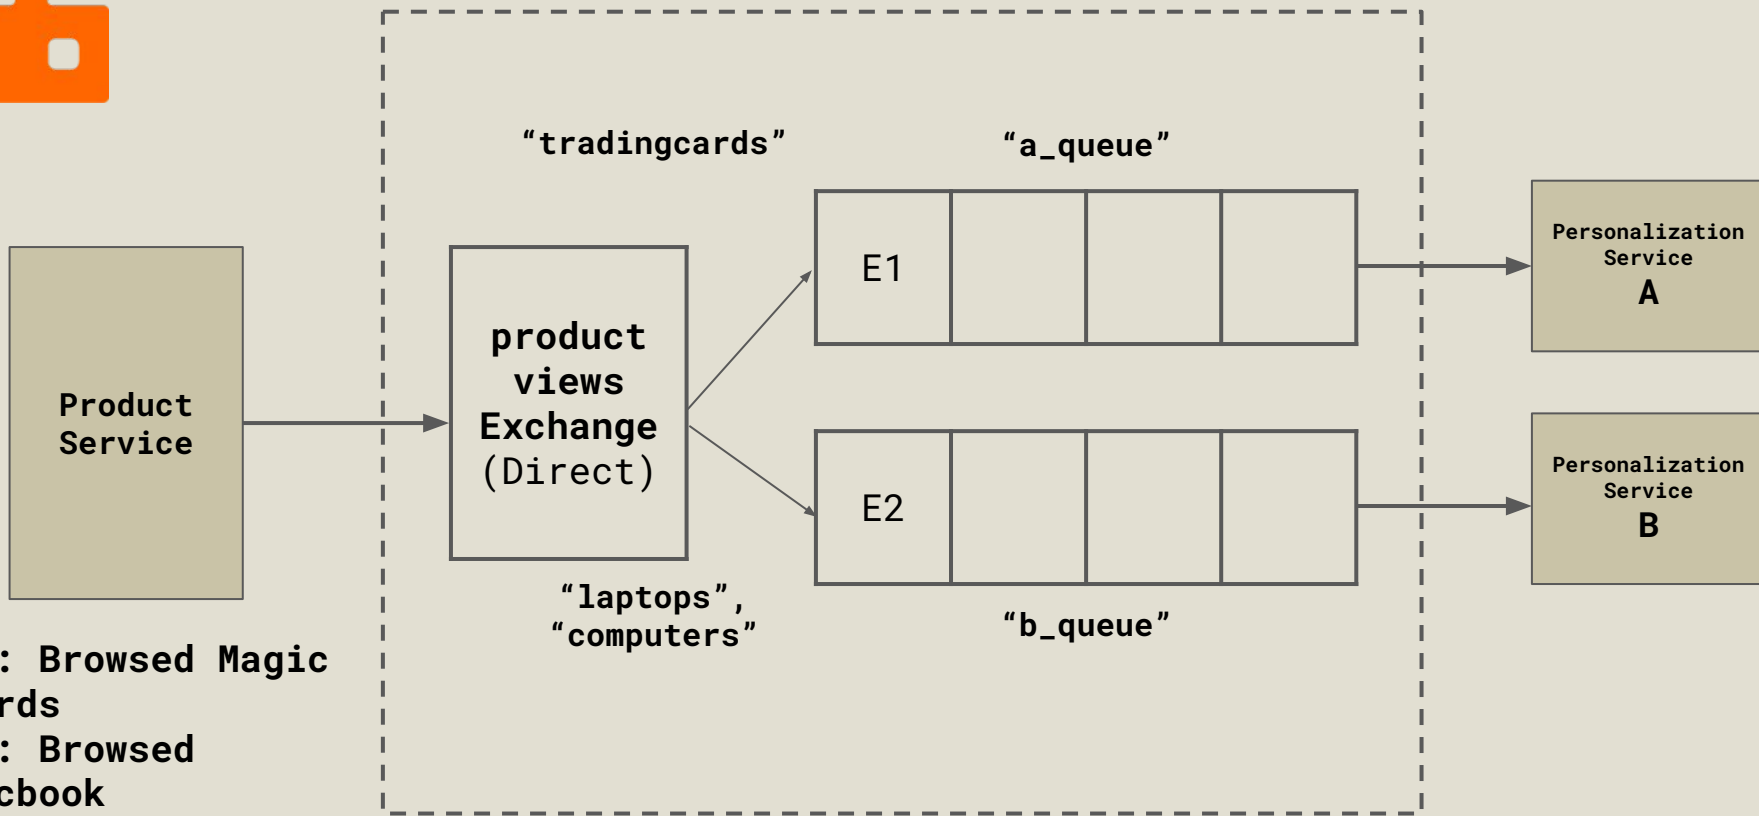
What if we want to handle more categories?



Personalization  
Service  
**B**

```
channel.queue_bind(exchange='productviews',  
                    queue='b_queue',  
                    routing_key='computers')
```

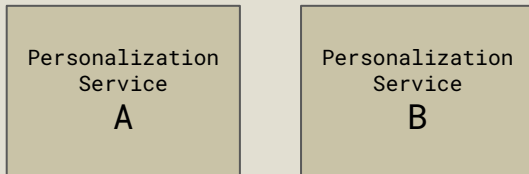
**Binding the same Queue to an Exchange**



E1: Browsed Magic Cards  
E2: Browsed Macbook

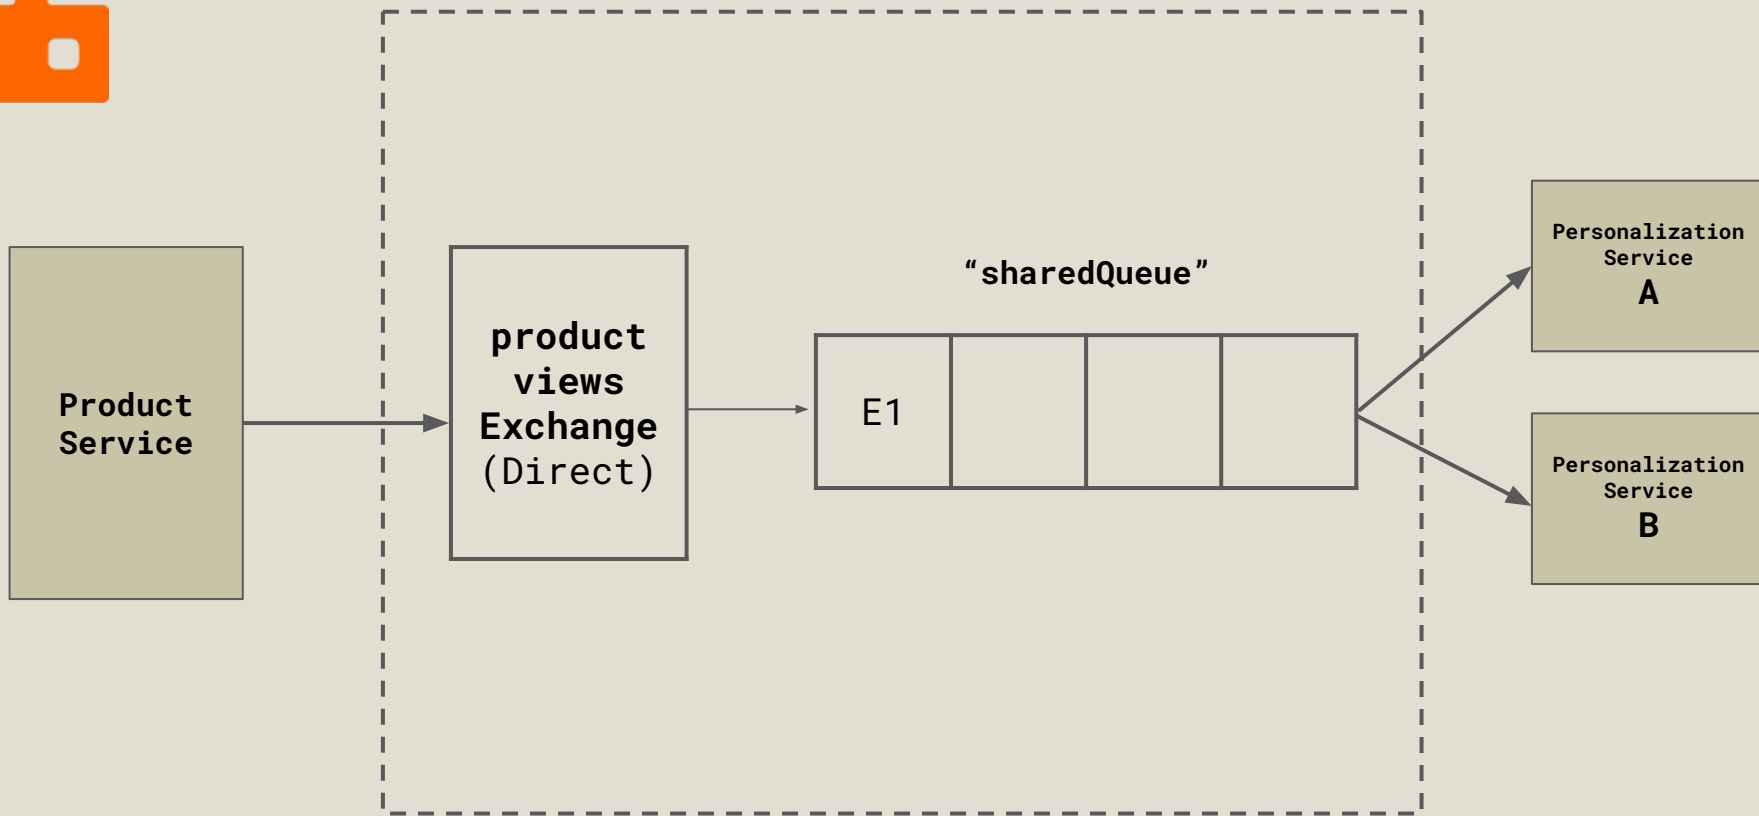
**"b\_queue" will now receive "computers" as well**



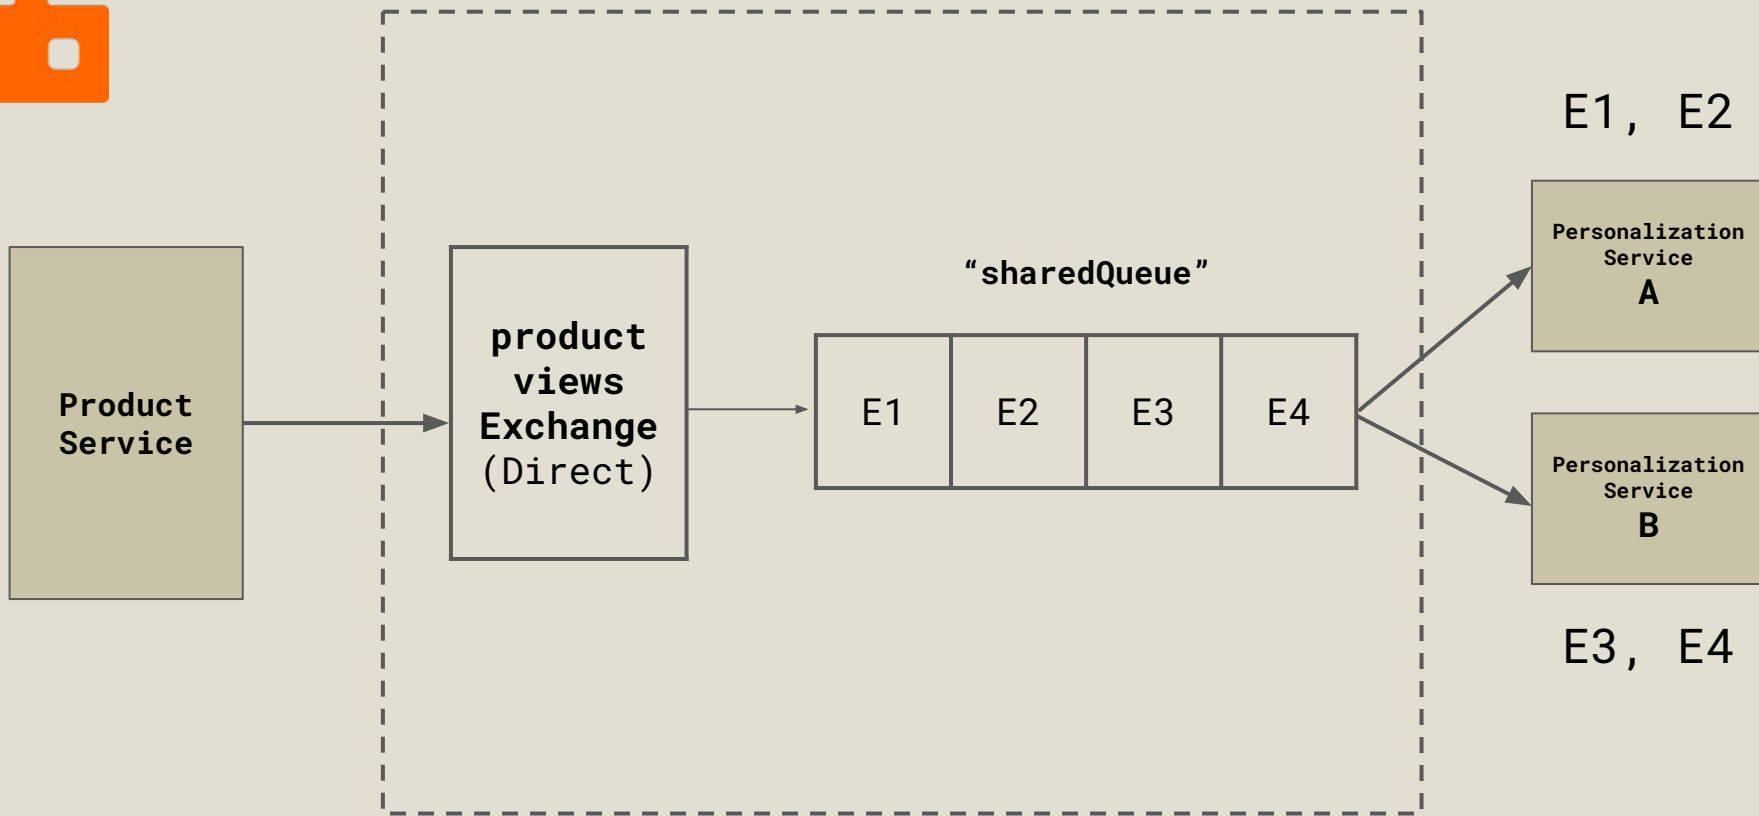


```
result = channel.queue_declare(queue='sharedQueue',  
                                exclusive=True)
```

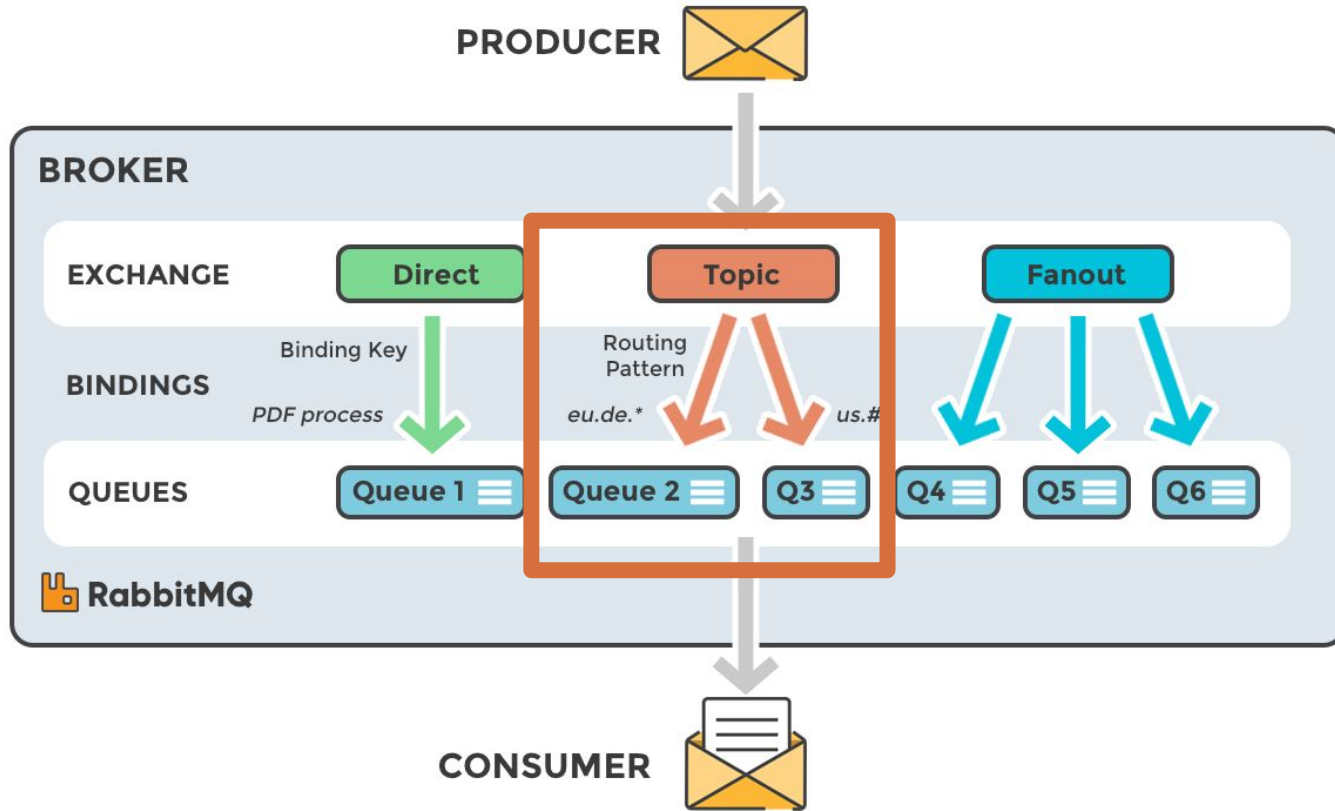
What if Services A and B bind the same queue name?



**Services can share a queue, but...**



**Events are load balanced (round-robin)**



Direct, Topic, and Fanout Exchanges



**E1: Any type of Magic Cards**



**Personalization  
Service  
A**

**E2: Apple Laptops**

**E3: Any types of PC**




**Personalization  
Service  
B**

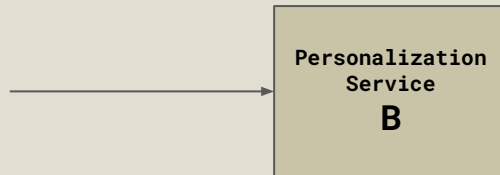
**Can we support more without many routing keys?**



E1: `category.cards.magic.tarkir`  
`category.cards.magic.*`



E2: `category.computers.laptops.apple`  
E3: `category.computers.pc.dell`



`category.computers.#`

## Topic Exchange: Use Wildcards



"The binding key must also be in the same form. The logic behind the topic exchange is similar to a direct one - a message sent with a particular routing key will be delivered to all the queues that are bound with a matching binding key.

However there are two important special cases for binding keys:

- \* (star) can substitute for exactly one word.**

- # (hash) can substitute for zero or more words."**


## **RabbitMQ: Topics**


<https://www.rabbitmq.com/tutorials/tutorial-five-python>




**\* (star) can substitute for exactly one word.**

**category.cards.magic.\***

category.cards.magic.tarkir 

category.cards.magic.innistrad 


category.cards.magic.innistrad.booster 

**# (hash) can substitute for zero or more words.**

**category.computers.#**

category.computers 

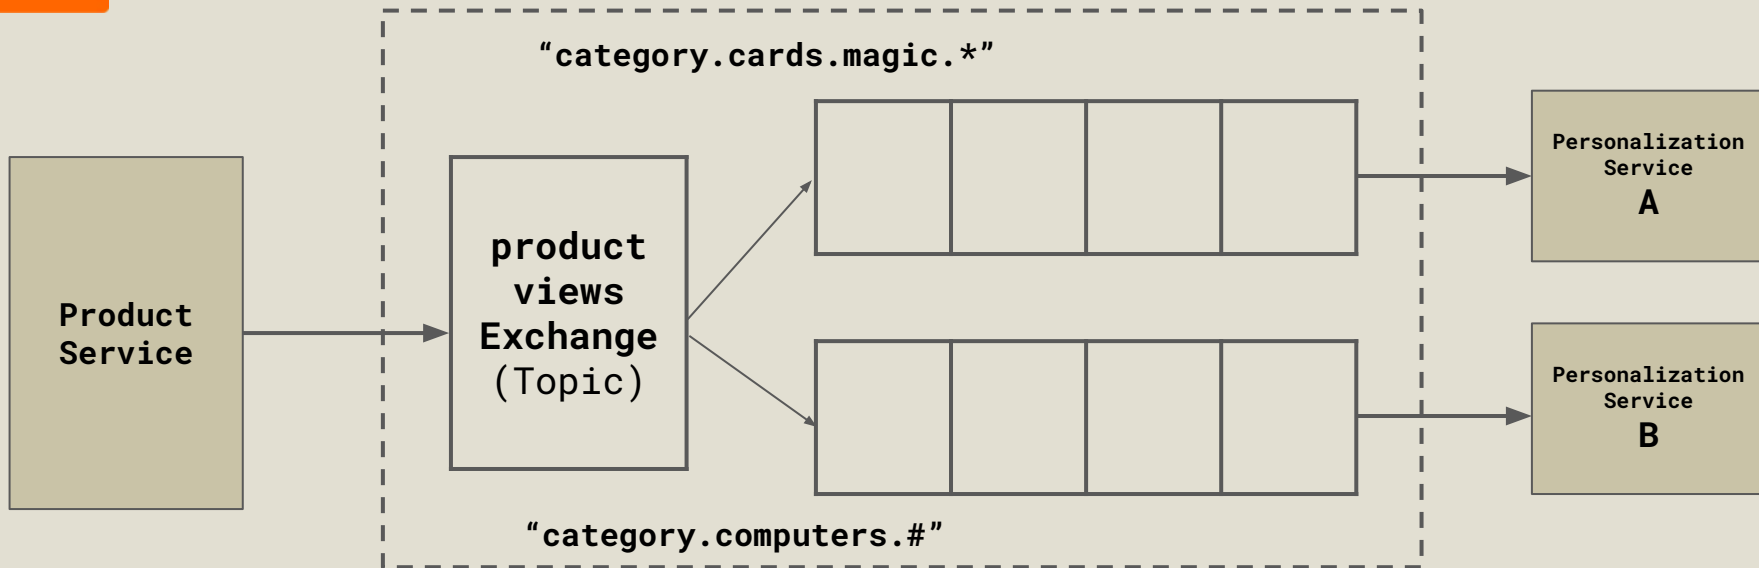
category.computers.macbook 

category.computers.macbook.pro 





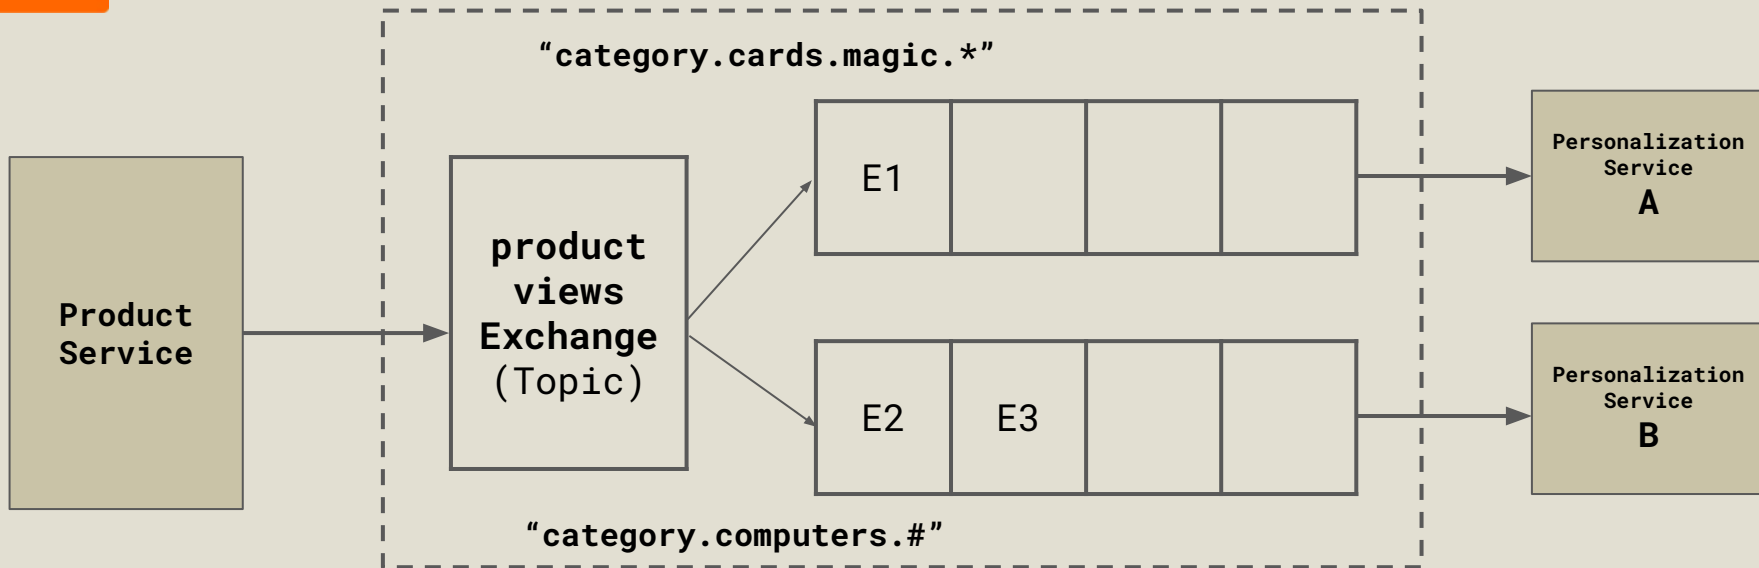
## Bind the queues to the Exchange via Routing Key



**Topic Exchange: allows wildcard bindings**



## Bind the queues to the Exchange via Routing Key



E1: `category.cards.magic.tarkir`  
E2: `category.computers.laptop.apple`  
E3: `category.computers.pc.dell`

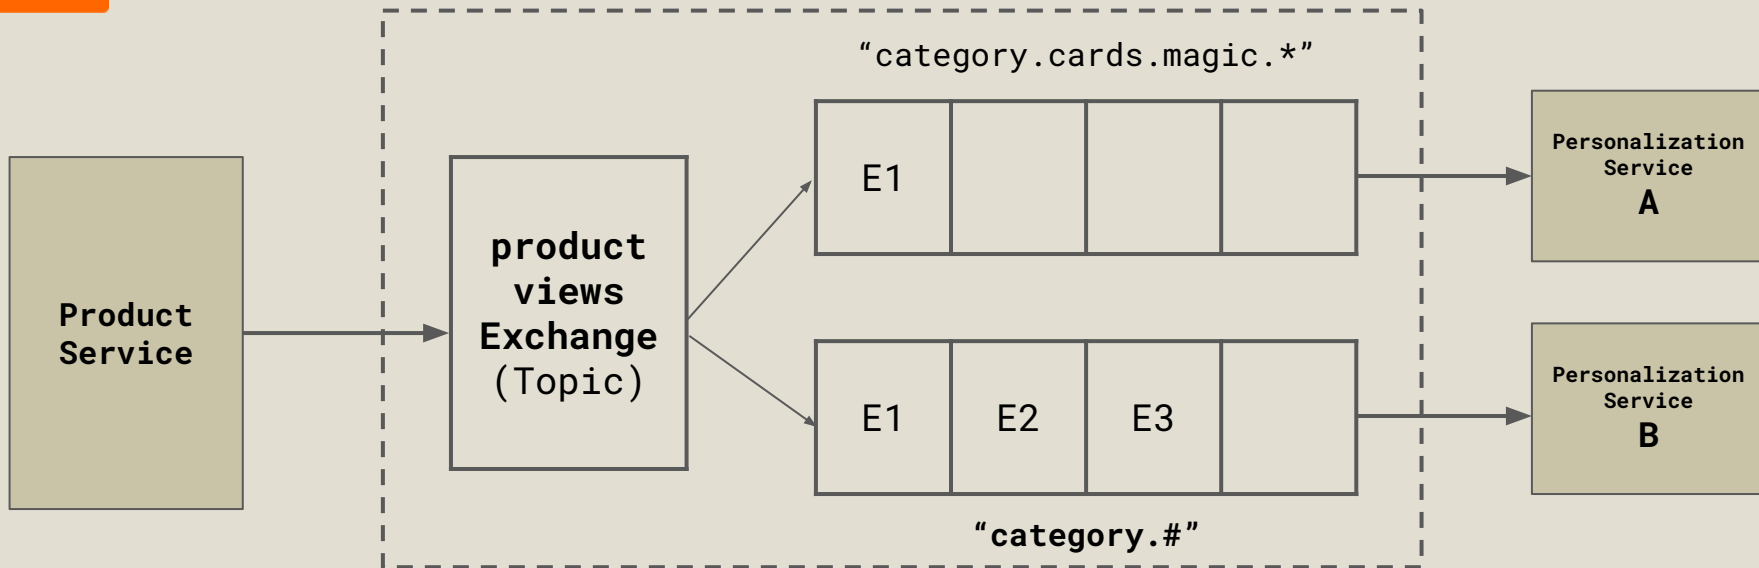
**Topic Exchange: allows wildcard bindings**



**What about "category.#"?**



## Bind the queues to the Exchange via Routing Key

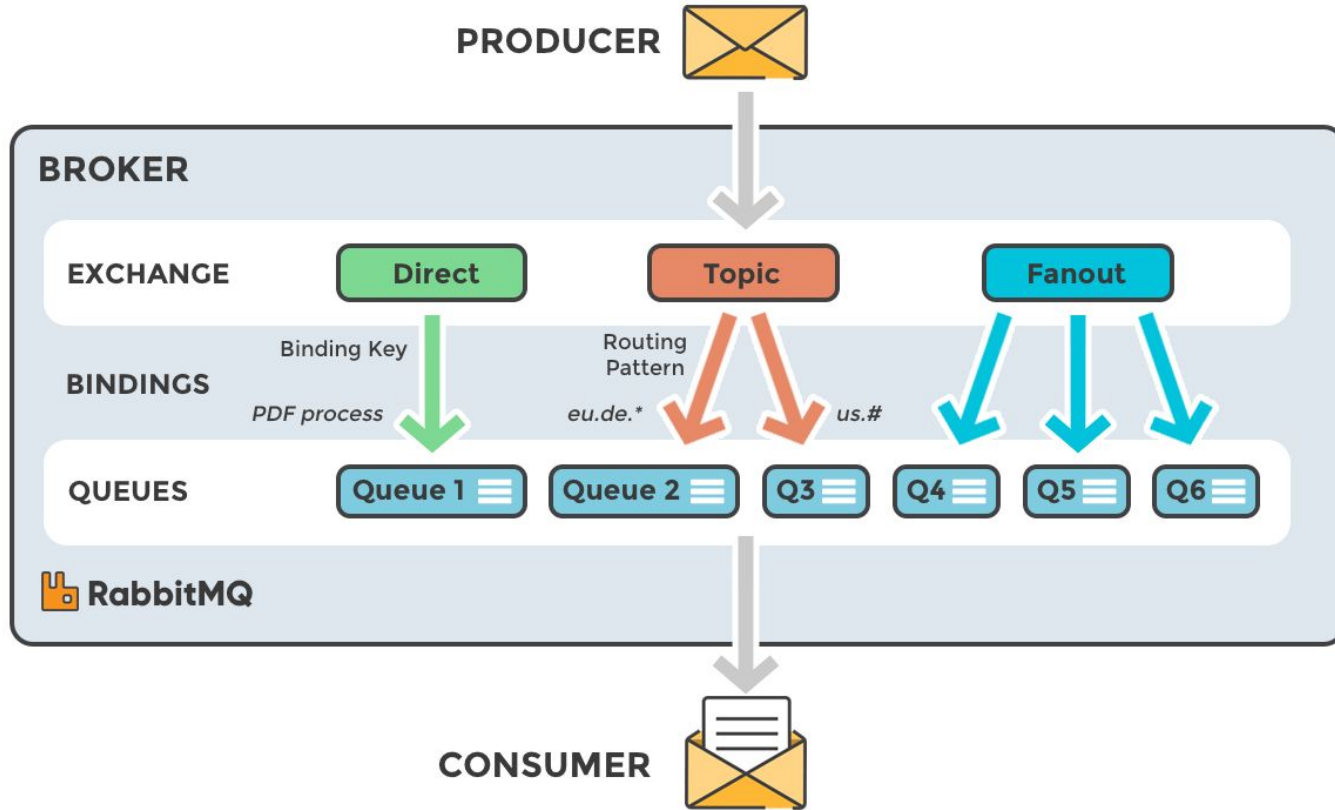


**E1: Browsed Magic Cards**

**E2: Browsed Macbook**

**E3: Browsed Dell PC**

## Topic Exchange: Wildcard bindings



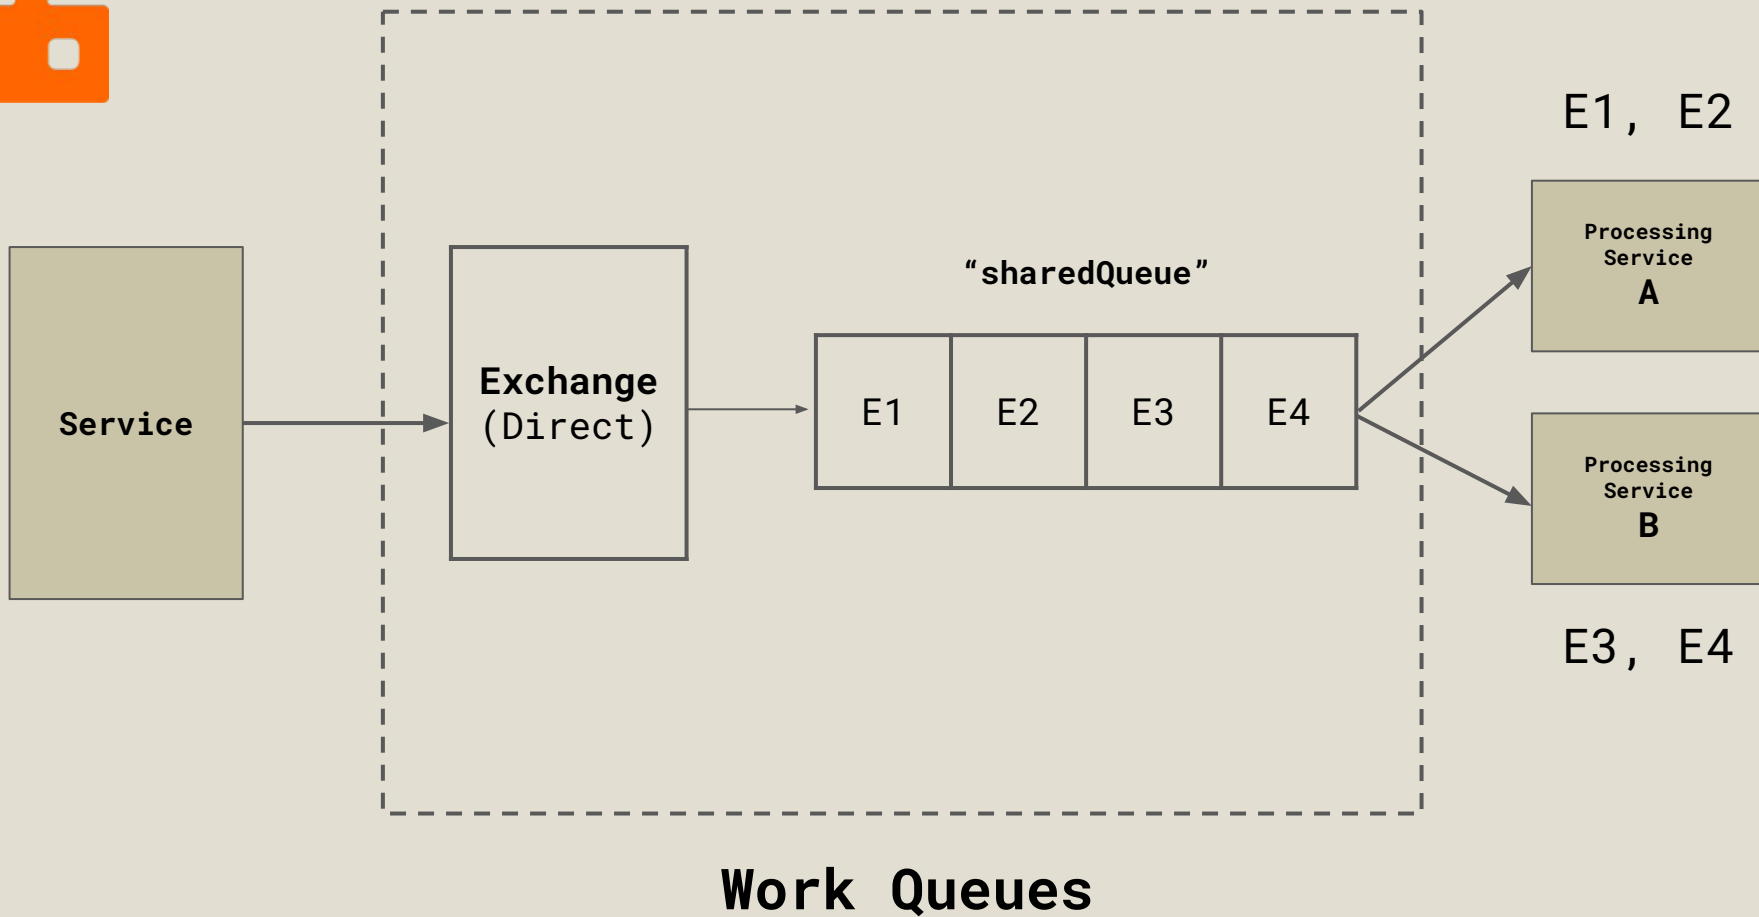
## Direct, Topic, and Fanout Exchanges



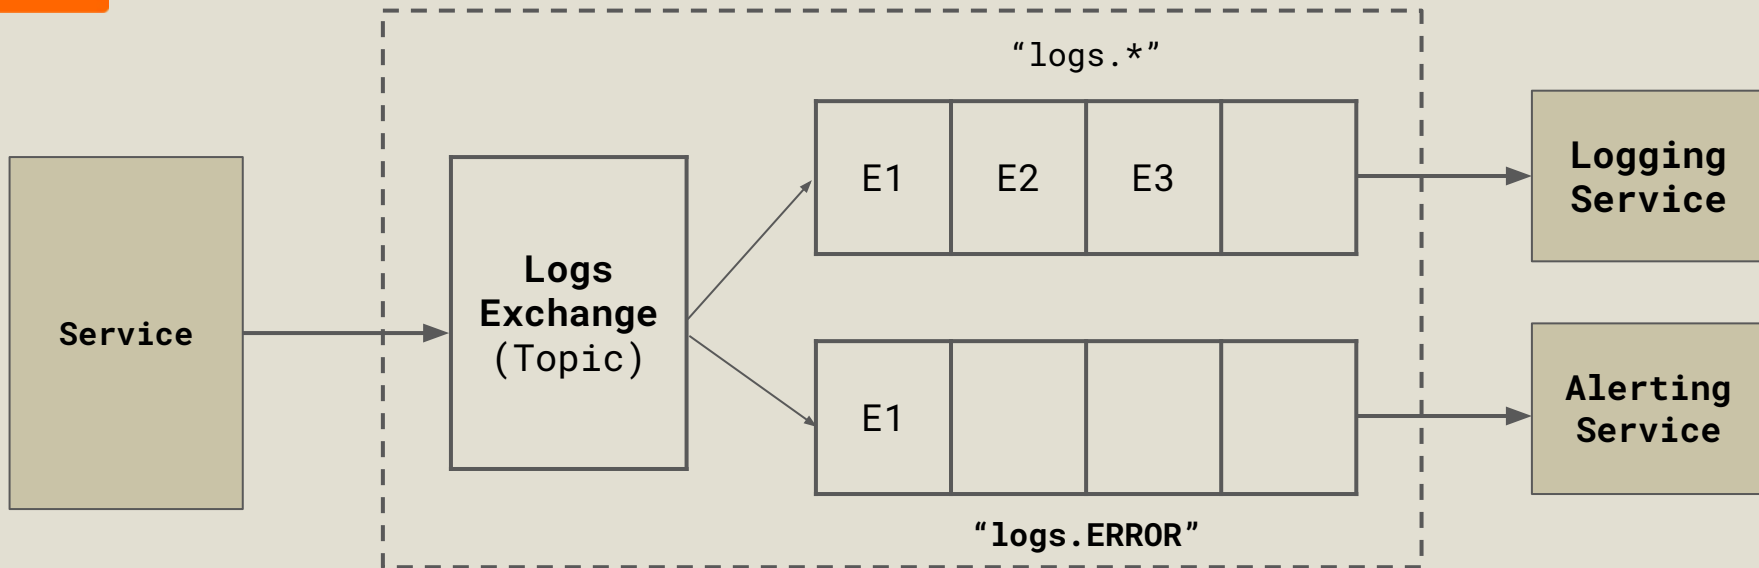
Exchange Type	Routing Key Supported	Use Cases
Fanout	N.A.	Receive everything from this exchange
Direct	"Routing Key" (no wildcards)	Receive events from selected keys without more fine-grained filters
Topic	"Routing Key" with wildcards, (* and #)	Allows fine-grained filters to only receive events of a subset of a routing key

## RabbitMQ: Exchange Binding Summary

# 5.1.2 other common patterns







E1: logs.ERROR  
E2: logs.INFO  
E3: logs.INFO

## Error Logging: Alert and Logging

Peer Discussion

Identify at least 3 more possible use cases for RabbitMQ queues in pairs.



Which exchange pattern should I use? (Direct, Fanout, or Topic)

~10 mins, pick 1 person to share after discussing.



# 5.1.3

## reliability

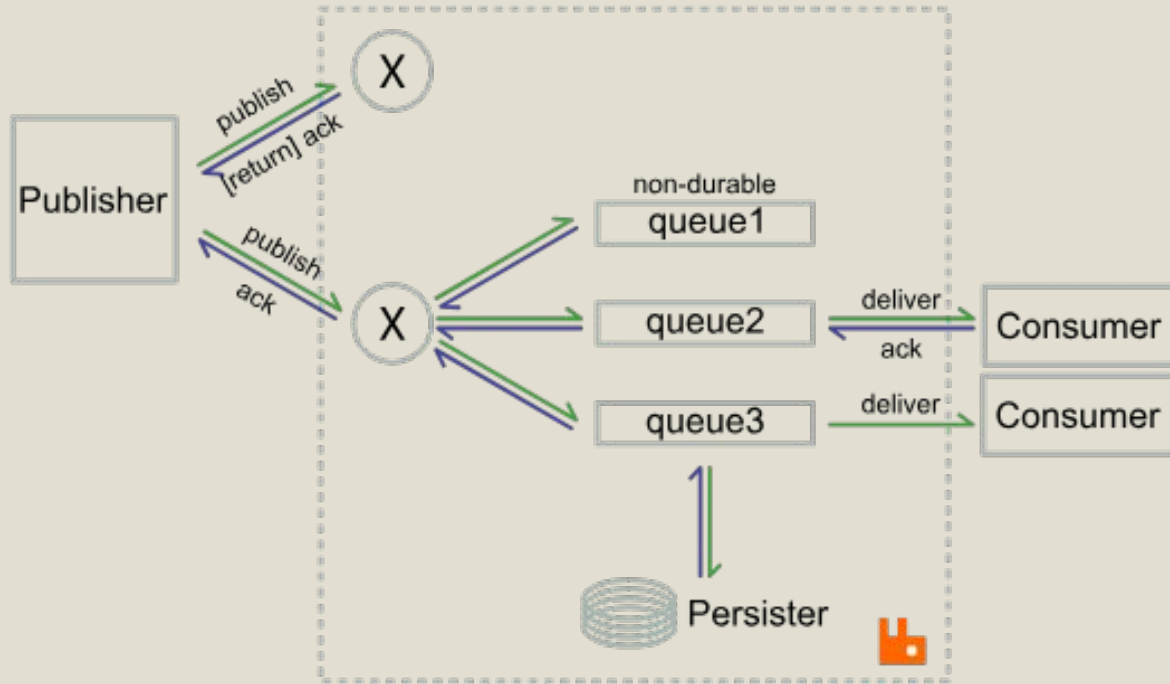
### options



**RabbitMQ Clustering**  
**Queues not replicated by default**



**RabbitMQ Clustering: Quorum Queues  
(Needs majority)**



## Publisher Confirms

<https://www.rabbitmq.com/blog/2011/02/10/introducing-publisher-confirms>



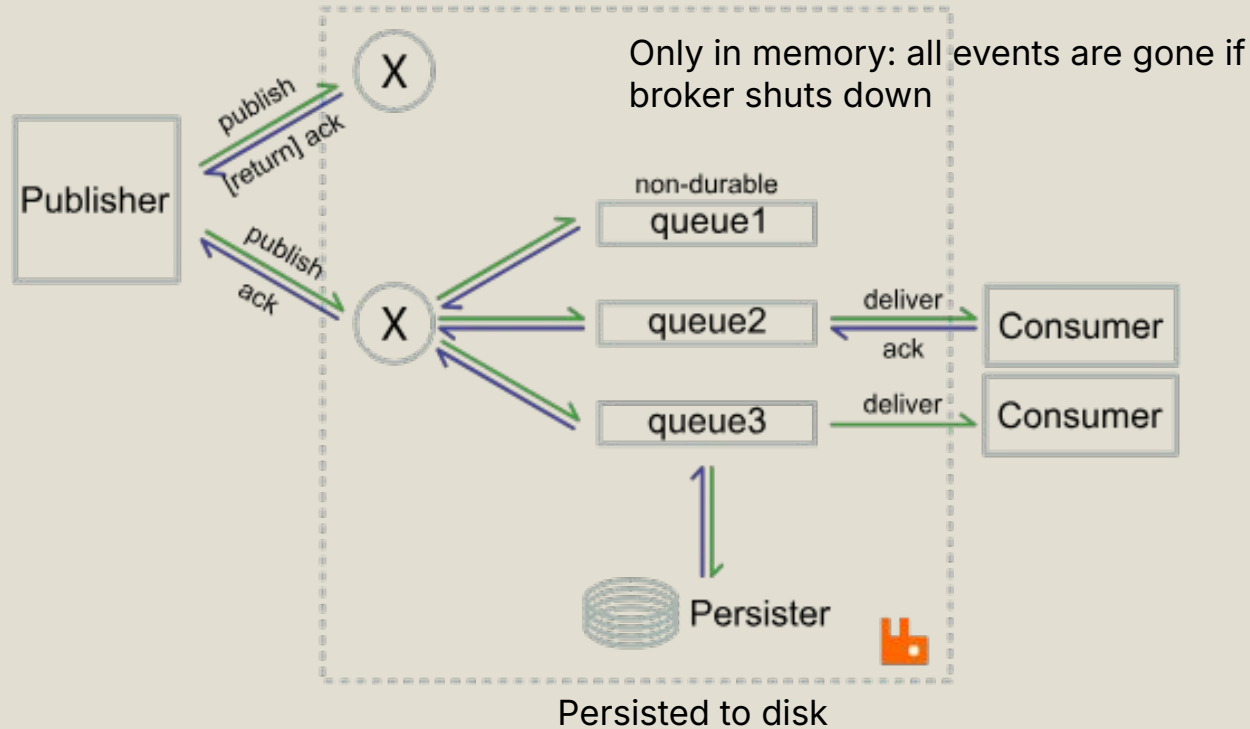
When the broker acknowledges a message, it assumes responsibility for it and informs the publisher that it has been handled successfully; what "handled successfully" means is context-dependent.

**The basic rules are as follows:**

- **an un-routable mandatory or immediate message is confirmed right after the *basic.return*;**
- **otherwise, a transient message is confirmed the moment it is enqueued; and,**
- **a persistent message is confirmed when it is persisted to disk or when it is consumed on every queue.**

## **Publisher Confirms**

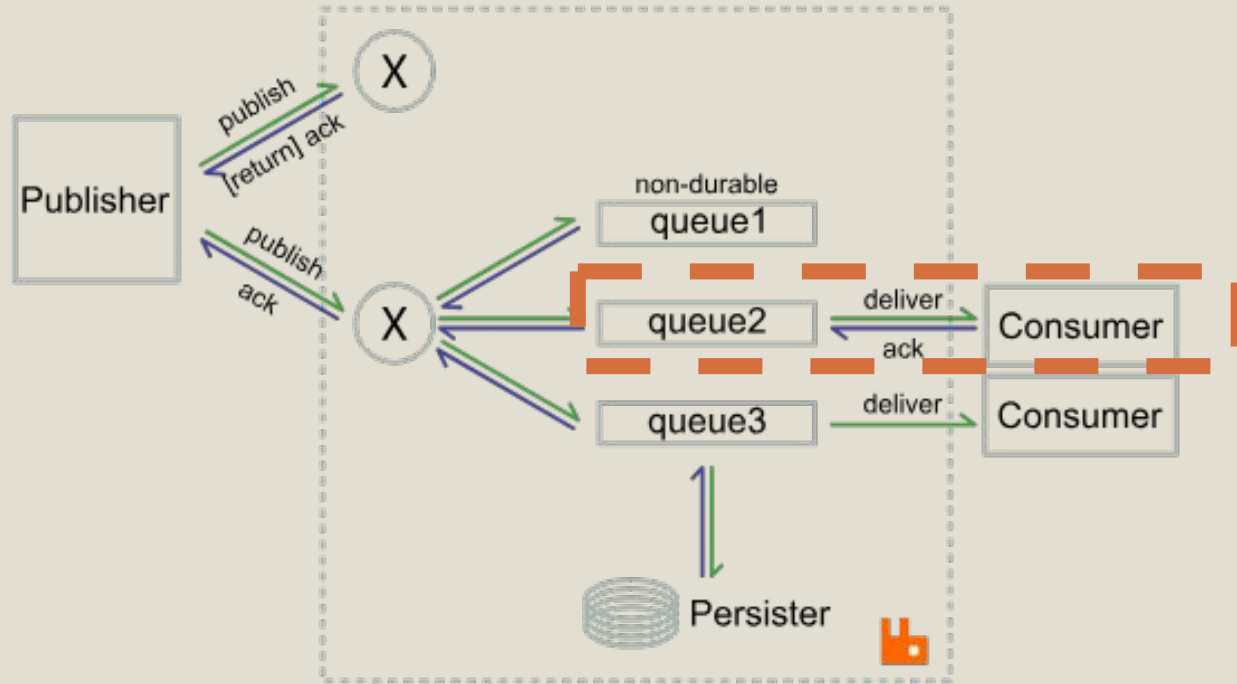
<https://www.rabbitmq.com/blog/2011/02/10/introducing-publisher-confirms>



## Publisher Confirms

<https://www.rabbitmq.com/blog/2011/02/10/introducing-publisher-confirms>





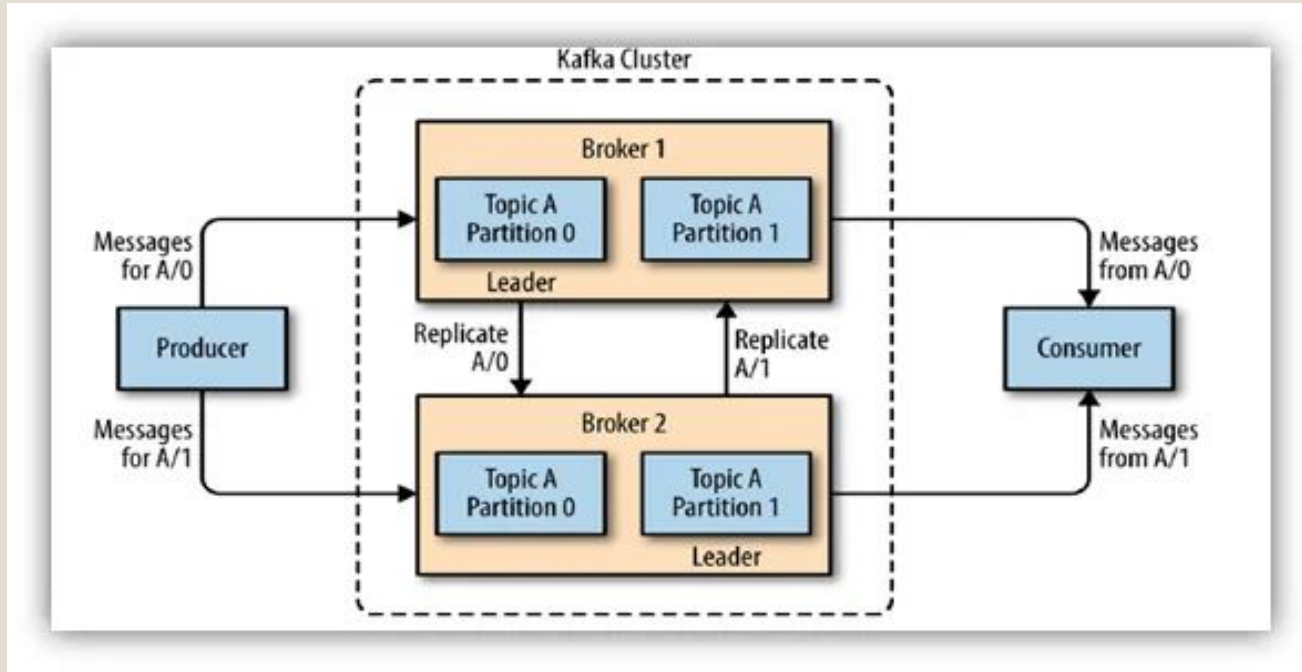
## Consumer Acknowledgements

<https://www.rabbitmq.com/docs/confirmations#consumer-acknowledgements>

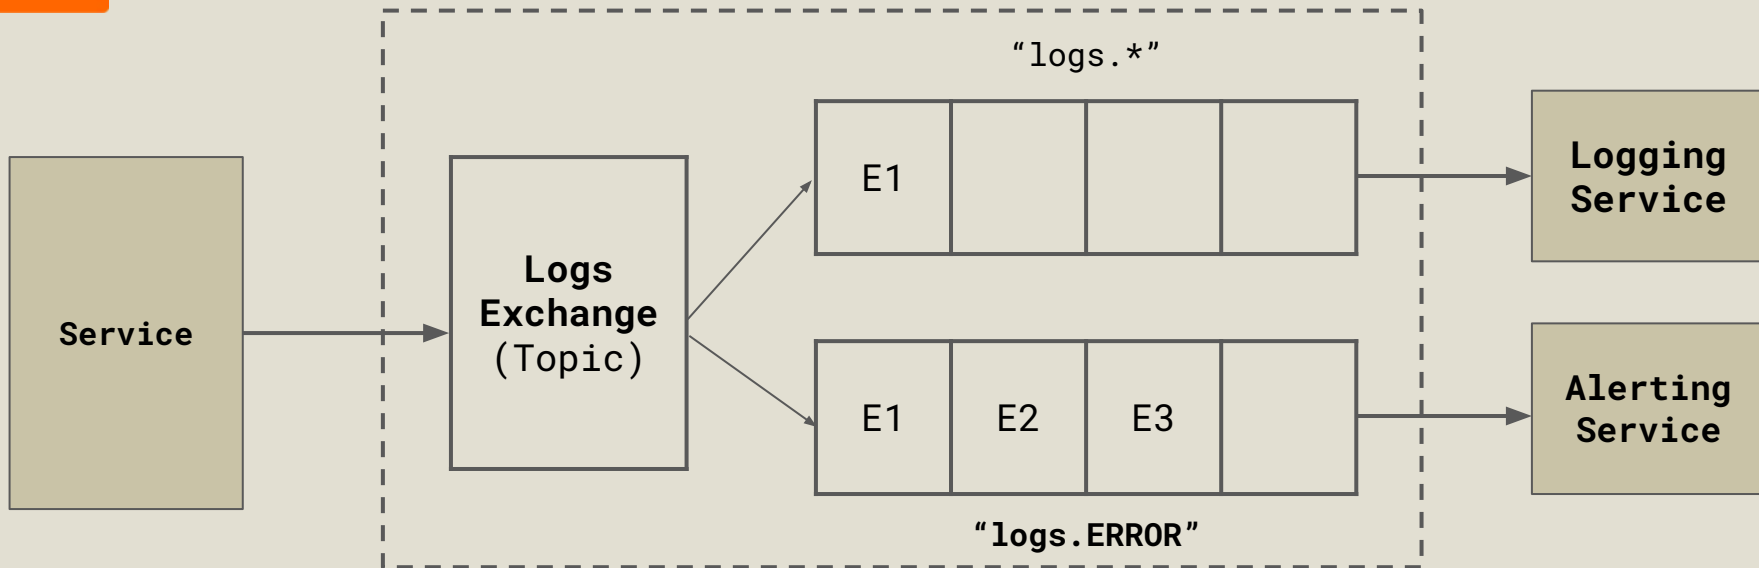
1. **Queues form the asynchronous data flow of distributed systems.** Data flows can be synchronous (eg. HTTP calls) or asynchronous.
2. **Flexible Implementation:** Choose between Topic, Fanout, or Direct exchanges to determine how you wish to publish and consume events to implement patterns you want (eg. Work Queues, filtered events, wildcards)
3. **Cluster or don't.** Clustering is optional for RabbitMQ, and can be used as a lightweight solution, or a full-featured reliable event system with clustering, quorum queues.

## **RabbitMQ: Summary**

## 5.2 kafka

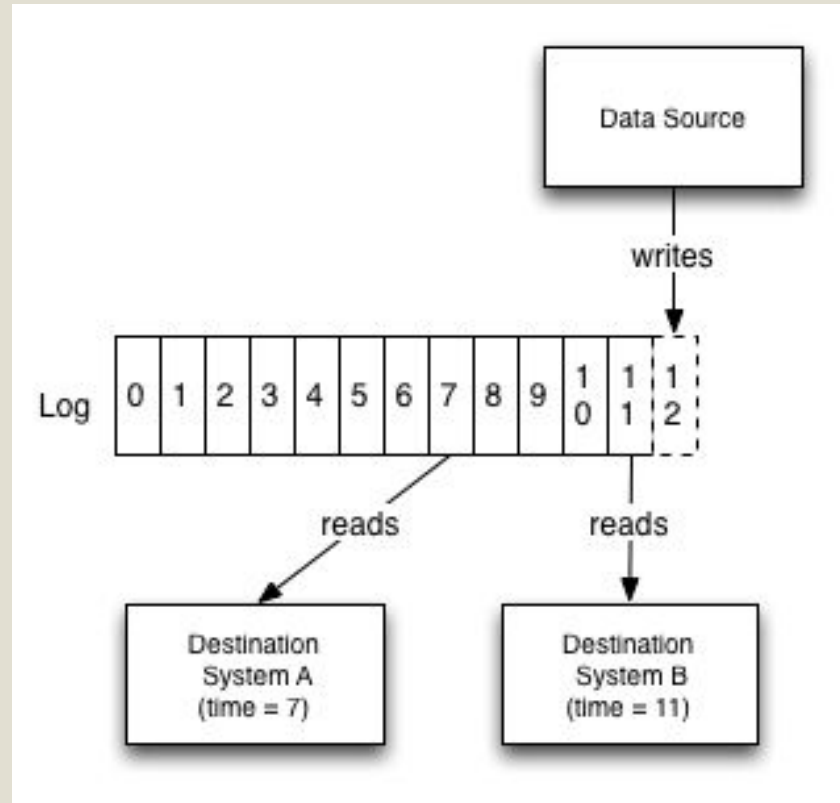


## Kafka: Architecture




E1: logs.ERROR  
E2: logs.INFO  
E3: logs.INFO

**RabbitMQ: Events Deleted After They are Consumed**




# Heart of Kafka: Append-Only Log

## Append-only log in Raft



**Figure 6:** Logs are composed of entries, which are numbered sequentially. Each entry contains the term in which it was created (the number in each box) and a command for the state machine. An entry is considered *committed* if it is safe for that entry to be applied to state machines.



# Raft

## Raft Algorithm: Quick Summary


<https://raft.github.io/raft.pdf>



"We used a few tricks in Kafka to support this kind of scale:

- Partitioning the log
- Optimizing throughput by batching reads and writes
- Avoiding needless data copies

In order to allow horizontal scaling we chop up our log into partitions:"



## Kafka: Architecture





Topic, eg.  
"product views"

Distributes partitions  
amongst consumers  
in the same group




Figure 1: Consumer Group

# Kafka: Topics and Consumer Group

<https://www.confluent.io/en-gb/blog/tutorial-getting-started-with-the-new-apache-kafka-0-9-consumer-client/>



Last "saved"  
checkpoint that Kafka  
knows this consumer  
has processed




Figure 2: The Consumer's Position in the Log

## Kafka: Topics and Consumer Group

<https://www.confluent.io/en-gb/blog/tutorial-getting-started-with-the-new-apache-kafka-0-9-consumer-client/>



"The high watermark is the offset of the last message that was successfully copied to all of the log's replicas."




Figure 2: The Consumer's Position in the Log

## Kafka: Topics and Consumer Group

<https://www.confluent.io/en-gb/blog/tutorial-getting-started-with-the-new-apache-kafka-0-9-consumer-client/>



"The log end offset is the offset of the last message written to the log."




Figure 2: The Consumer's Position in the Log

## Kafka: Topics and Consumer Group

<https://www.confluent.io/en-gb/blog/tutorial-getting-started-with-the-new-apache-kafka-0-9-consumer-client/>



Exchange Type	Routing Key Supported	Use Cases
Fanout	N.A.	Receive everything from this exchange
Direct	"Routing Key" (no wildcards)	Receive events from selected keys without more fine-grained filters
Topic	"Routing Key" with wildcards, (* and #)	Allows fine-grained filters to only receive events of a subset of a routing key

**Can we do the same as RabbitMQ? Hint: No**




"If the producer specifies a key, Kafka applies a hash function to the key, which results in a numerical value.

This hash value is then used to determine which partition the message will be sent to.

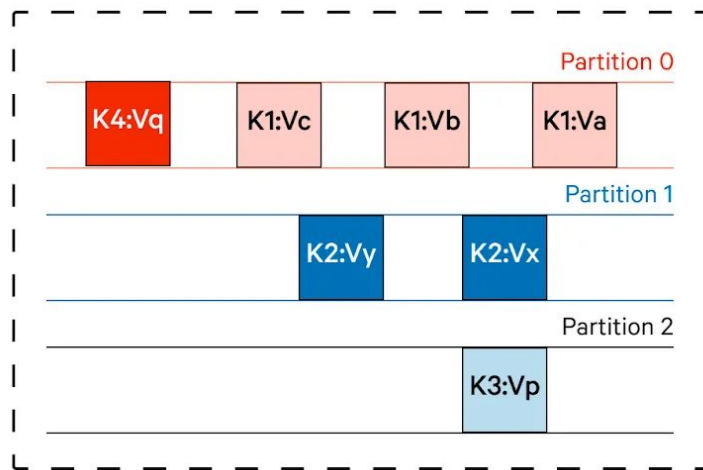
**In simpler terms, Kafka uses the key to ensure that all messages with the same key are sent to the same partition, allowing for grouping of related messages."**

## **Kafka: Partitioning by Message Key**

<https://www.confluent.io/learn/kafka-message-key/>




Messages to be published  
(in chronological order, K1:Va being the first)



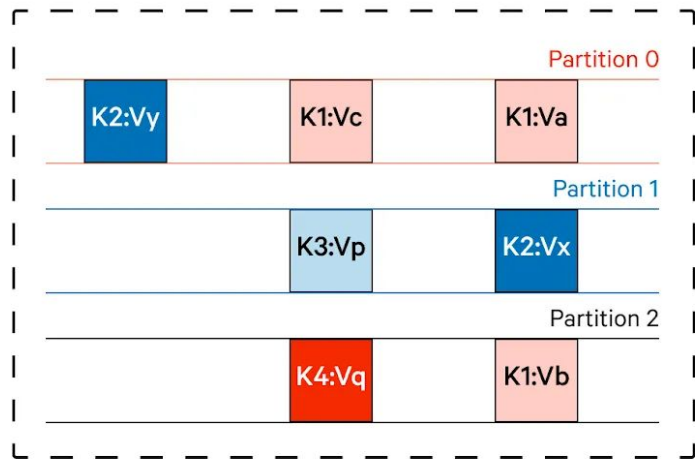
Default partitioner - Partition is decided  
on the basis of key hash

# Kafka: Default Partitioner

<https://www.redpanda.com/guides/kafka-tutorial-kafka-partition-strategy>



Messages to be published  
(in chronological order, K1:Va being the first)



Messages are sent to partitions in a  
round robin manner


# Kafka: Round Robin Partitioner

<https://www.redpanda.com/guides/kafka-tutorial-kafka-partition-strategy>

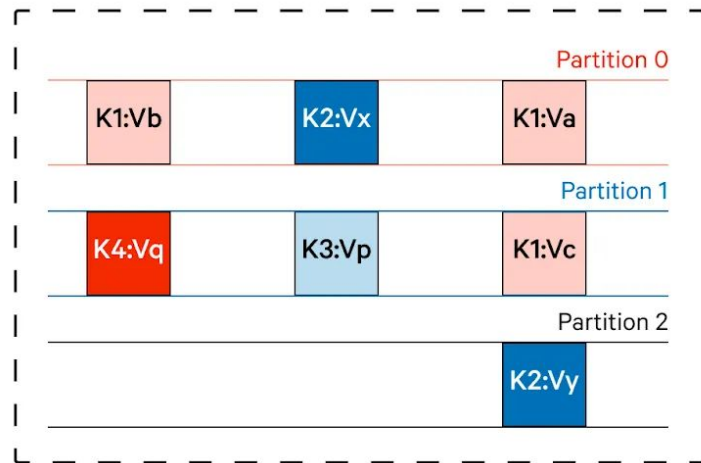




Similar to round-robin, but in batches



Messages to be published  
(in chronological order, K1:Va being the first)



Sticky partition changes only when batch is full. Also notice that records with same key are not guaranteed to be in the same partition

## Kafka: Uniform Sticky Partitioner

<https://www.redpanda.com/guides/kafka-tutorial-kafka-partition-strategy>



```
public class PowerUserPartitioner implements Partitioner {

    public void configure(Map<String, ?> configs) {}

    public int partition(String topic, Object key, byte[] keyBytes,
        Object value, byte[] valueBytes, Cluster cluster) {

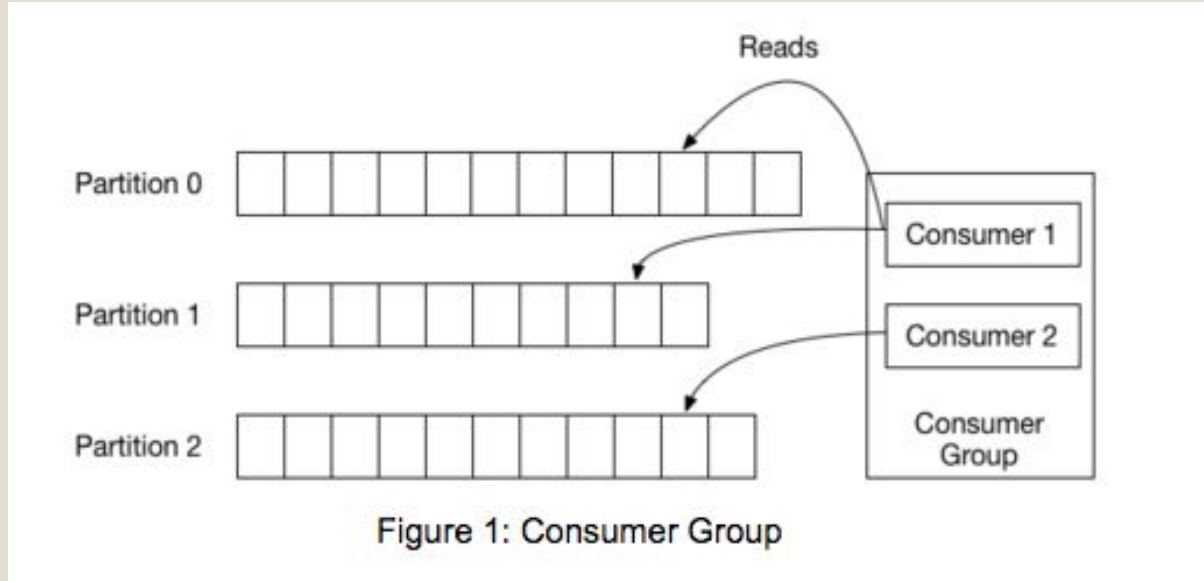
        List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
        int numPartitions = partitions.size();
        if ((keyBytes == null) || (!(key instanceof String)))
            throw new InvalidRecordException("Record must have a valid string key");

        if (((String) key).equals("CEO"))
            return numPartitions - 1; // Messages with key "CEO" will always go to the last
partition


        // Other records will get hashed to the rest of the partitions
        return Math.abs(Utils.murmur2(keyBytes)) % (numPartitions - 1);
    }
}
```

## Kafka: Custom Partitioner

<https://www.redpanda.com/guides/kafka-tutorial-kafka-partition-strategy>



**Kafka: Suitable for Large Scale Fanout  
at High Throughput**



**Fanout: More Troublesome with RabbitMQ**  
**Need to bind queues to consumers**



Version: 4.0

# Streams and Super Streams (Partitioned Streams)

## What is a Stream

RabbitMQ Streams is a persistent replicated data structure that can complete the same tasks as queues: they buffer messages from producers that are read by consumers. However, streams differ from queues in two important ways: how messages are stored and consumed.

Streams model an append-only log of messages that can be repeatedly read until they expire. Streams are always persistent and replicated. A more technical description of this stream behavior is “non-destructive consumer semantics”.

To read messages from a stream in RabbitMQ, one or more consumers subscribe to it and read the same messages as many times as they want.

## RabbitMQ: Streams offer something similar

# 5.2.1

## reliability

### options




Figure 2: The Consumer's Position in the Log

**Consumer crashes: what happens?**




Figure 1: Consumer Group

**Kafka reassigns another consumer  
in the same group to take over**






Figure 2: The Consumer's Position in the Log

**New Consumer restarts from offset 1**


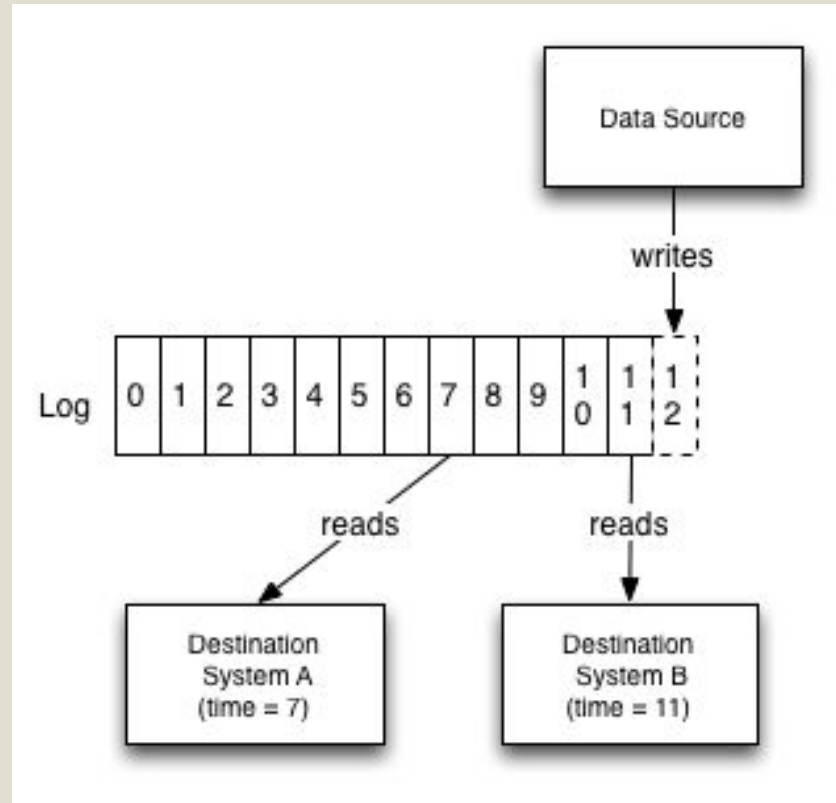


Figure 2: The Consumer's Position in the Log

**Can process until high watermark point**

`auto.commit.interval.ms`. The  
default interval is 5 seconds.

**Consumer offset default: Auto Commit 5 seconds**



## Publisher Acknowledgement





**acks=0** (No acknowledgment): The producer doesn't wait for any acknowledgment. It sends the message and doesn't care if it's received.

**acks=1** (Leader acknowledgment): The producer waits for an acknowledgment from the leader broker of the topic partition where the message is sent. This means the message is considered sent once the leader broker confirms receipt.

**acks=all** (Replica acknowledgment): The producer waits for acknowledgment from all in-sync replicas of the partition. This provides the highest level of durability but can introduce more latency.

## Publisher Acknowledgement: Levels

<https://medium.com/@monish.koppa/message-acknowledgement-in-apache-kafka-8daa8524c8e7>



Range Replication in CockroachDB

**Number of Nodes** and **Number of Replicas (Replication Factor)** available determines the number of node failures we can tolerate.

Number of replicas = 3 by default in CockroachDB

## Range Replication on Nodes

<https://www.cockroachlabs.com/blog/the-new-stack-meet-cockroachdb-the-resilient-sql-database/>



"Apache Kafka replicates the event log for each topic's partitions across a configurable number of servers.

**This replication factor is configured at the topic level, and the unit of replication is the topic partition.**

This enables automatic failover to these replicas when a server in the cluster fails so messages remain available."

## Kafka: Replication Factor

<https://kafka.apache.org/documentation/#replication>



Topics are added and modified using the topic tool:

```
$ bin/kafka-topics.sh --bootstrap-server localhost:9092  
--create --topic my_topic_name \  
    --partitions 20 --replication-factor 3 --config x=y
```

## Kafka: Replication Factor

<https://kafka.apache.org/documentation/#replication>




1. **Kafka's key data structure is the append-only log.** Append-only logs allow for high-throughput fanout event systems to be built easily.
2. **Different Partitioners provide flexibility for message routing to different partitions.** If required, you can write a custom partitioner to give you the ability to control how you wish to balance your partitions.
3. **Consumer groups and offset tracking provide flexibility for consumers to stop and resume.** Kafka can easily reassign consumers to replace dead consumers and resume events processing.

## **Kafka : Summary**

**p.s.**

**Current Trends**



```
LISTEN virtual;
```

```
NOTIFY virtual;
```

Asynchronous notification "virtual" received from server process with PID 8448.

```
NOTIFY virtual, 'This is the payload';
```

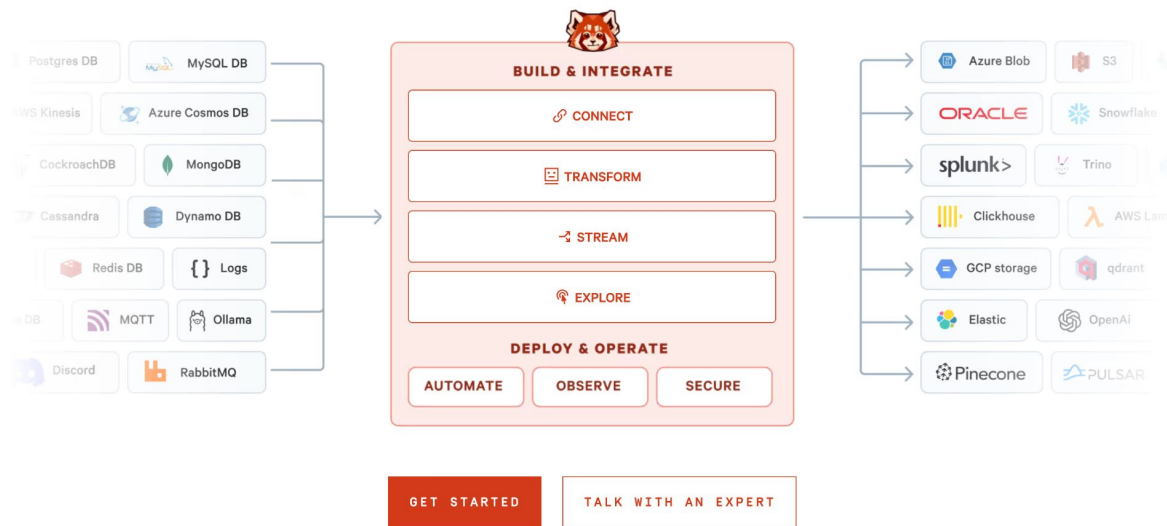
Asynchronous notification "virtual" with payload "This is the payload" received from server process with PID 8448.

## PostgreSQL for Events

<https://www.postgresql.org/docs/current/sql-notify.html>

# All-in-one platform for streaming, AI, and beyond

Redpanda is a complete Kafka® compatible streaming data platform with built-in developer tools and a growing ecosystem of connectors that's easy to integrate, simple to scale, and secure to run in any environment.



Spin up within minutes | No credit card required

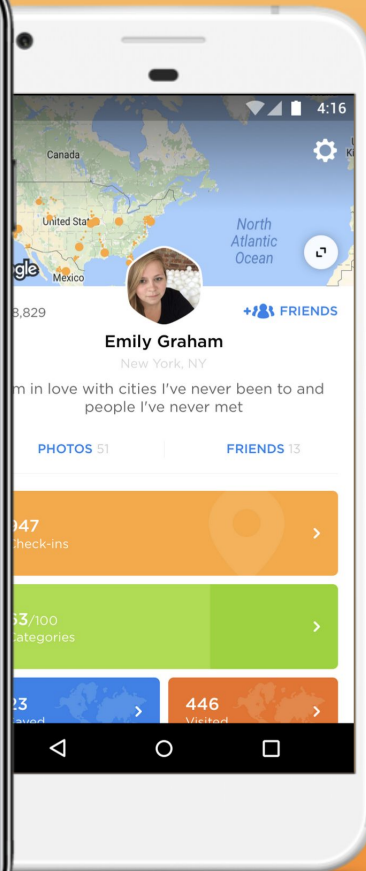
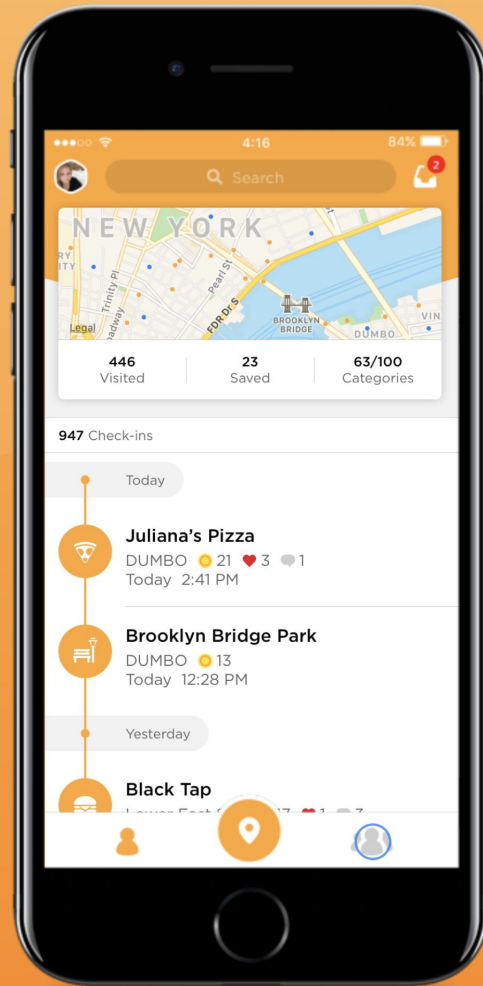
## Redpanda: Kafka Compatible

<https://www.redpanda.com/>

# Assignment

## Extend a Distributed Social Media Platform





## Assignment

- W1 → Distributed Social Media Research
- W2 → Review W1 + Build PoC
- W3 → Design 1st Draft of System + Sharing
- W4 → Implementation
- W5 → Implementation (Queues)**
- W6 → Implementation (Load Testing)
- W7 → Technical Retrospective
- W8 → Complete System Implementation

#### Assignment #4

- Integrate an event / queue system into your system for at least one or more of these use cases:
  - Handle and process incoming Posts or events, eg. Firehose or Jetstream for AT Protocol
  - Process outgoing Posts or events, eg. Outbox for ActivityPub
  - Any other use cases you can think of
- Share your implementation in code, and share a link to your project (eg. GitHub or Gitlab)